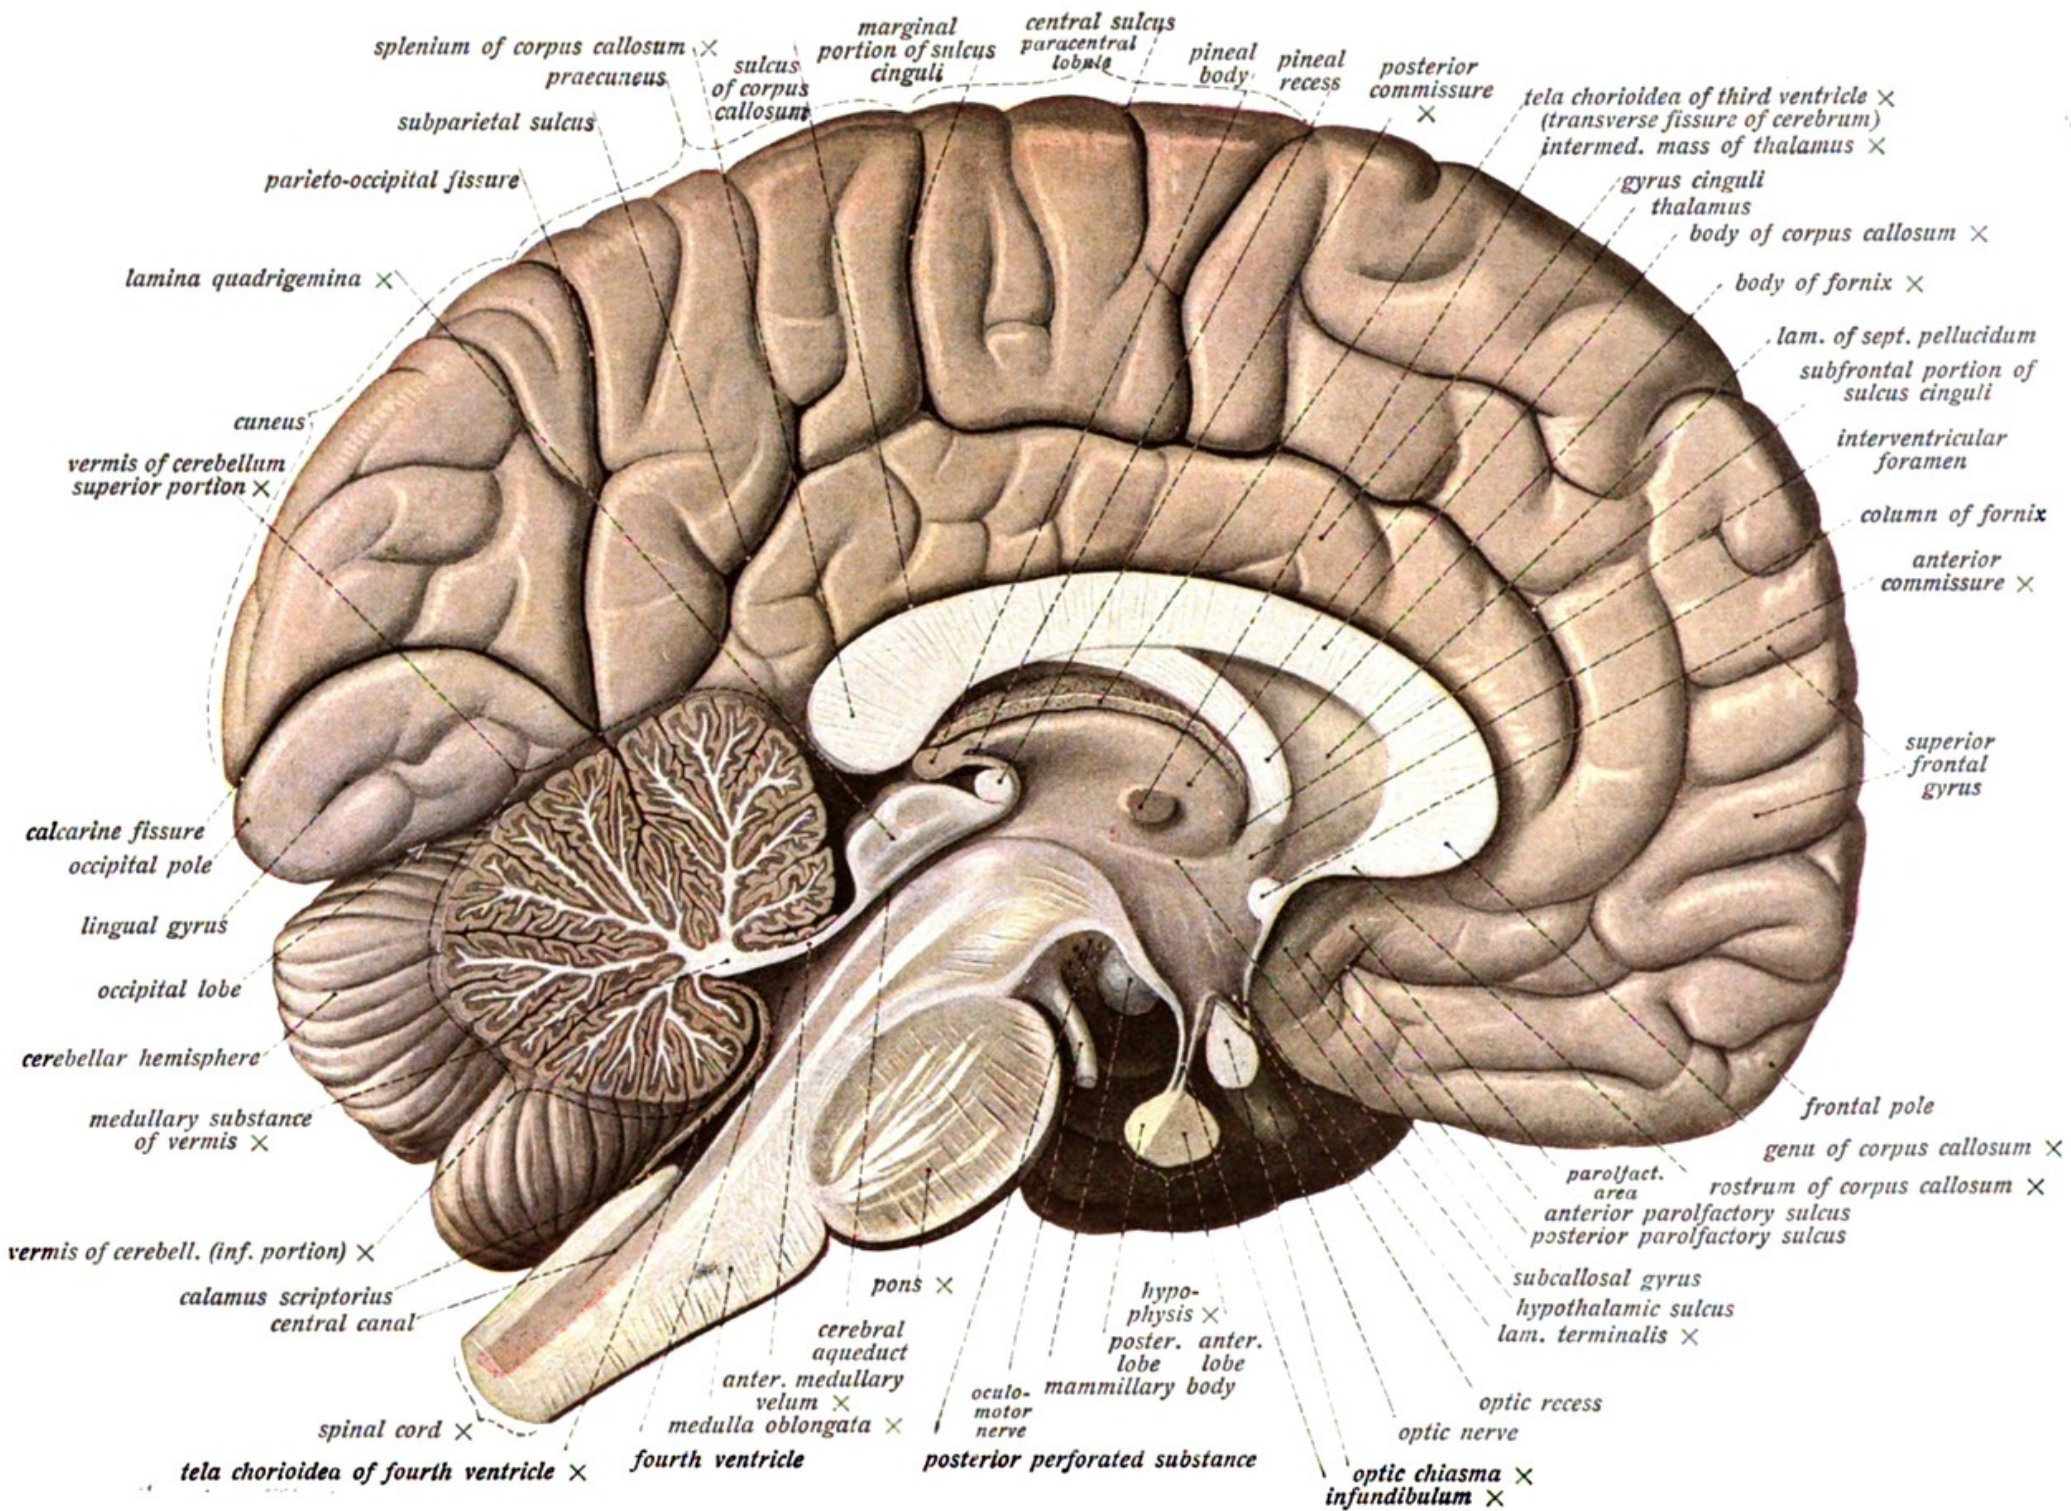


P08820 Introduction to Machine Learning

Week 5: Neural Networks

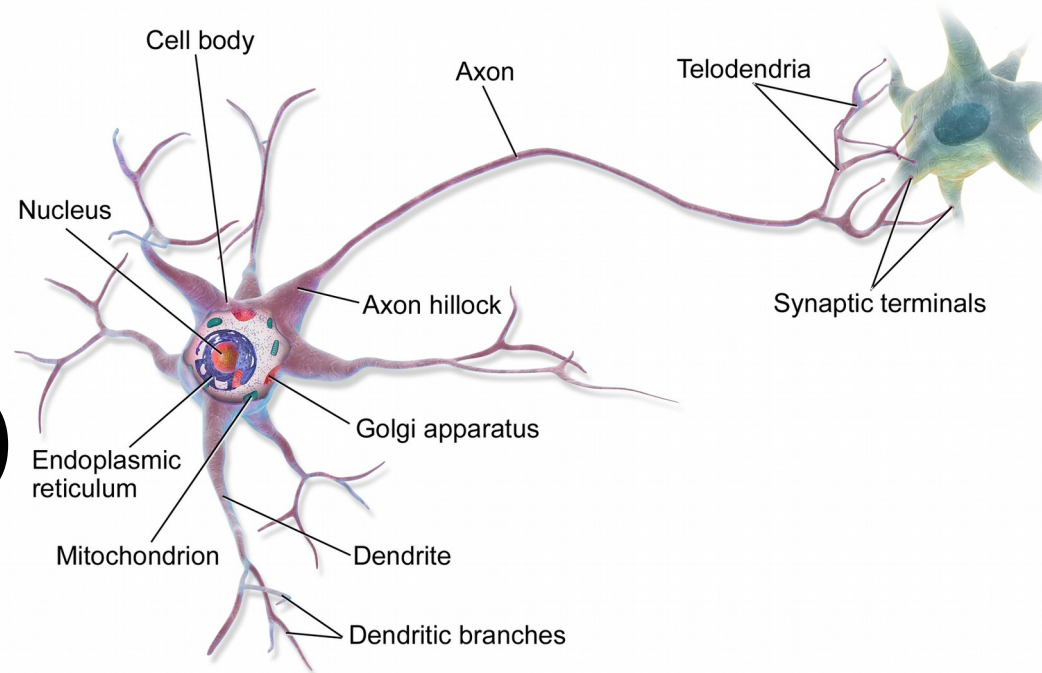
Matthias Rolf

mrolf@brookes.ac.uk



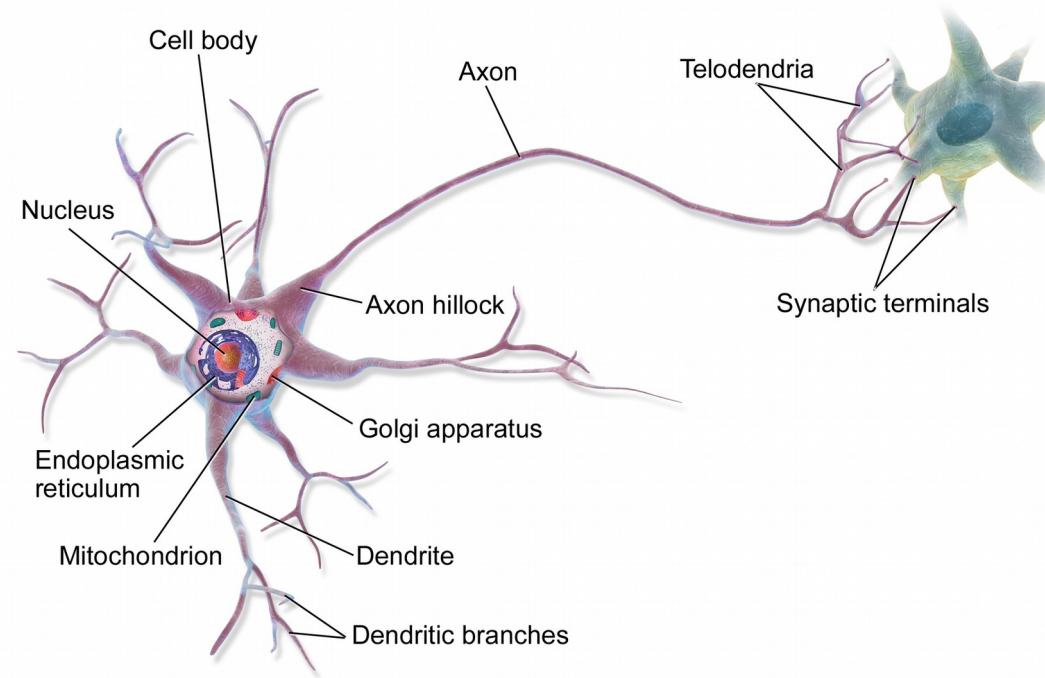
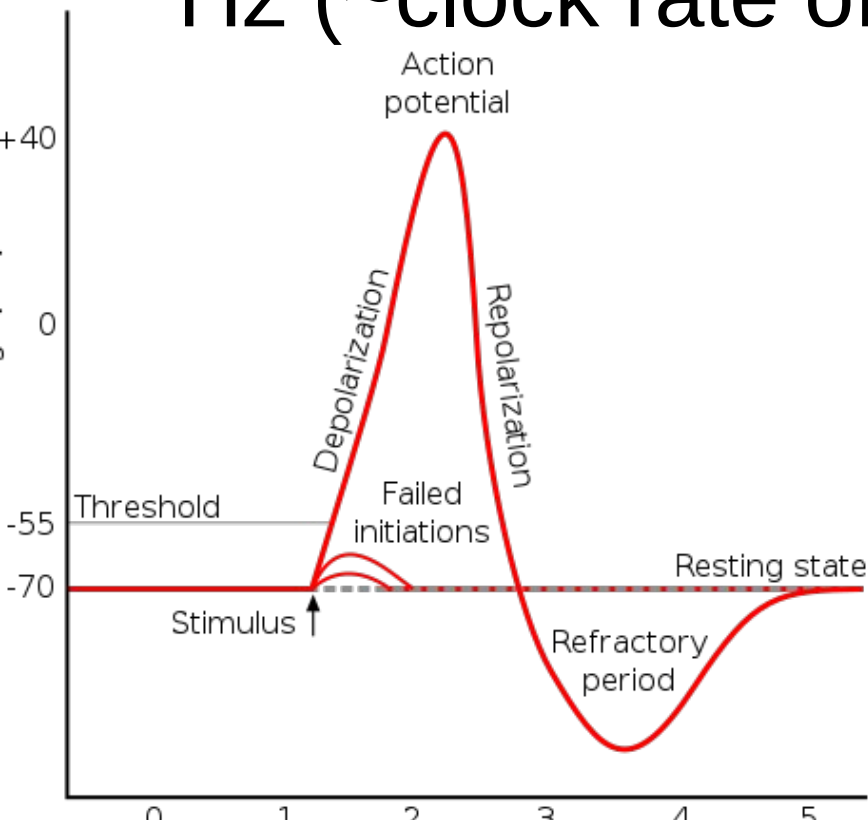
Neurons

- Relatively simple information processing units – but ca 100 billion
- Collect electric charge through dendrites
- Release spike of electric charge as a result through axon (“nerve fiber”)



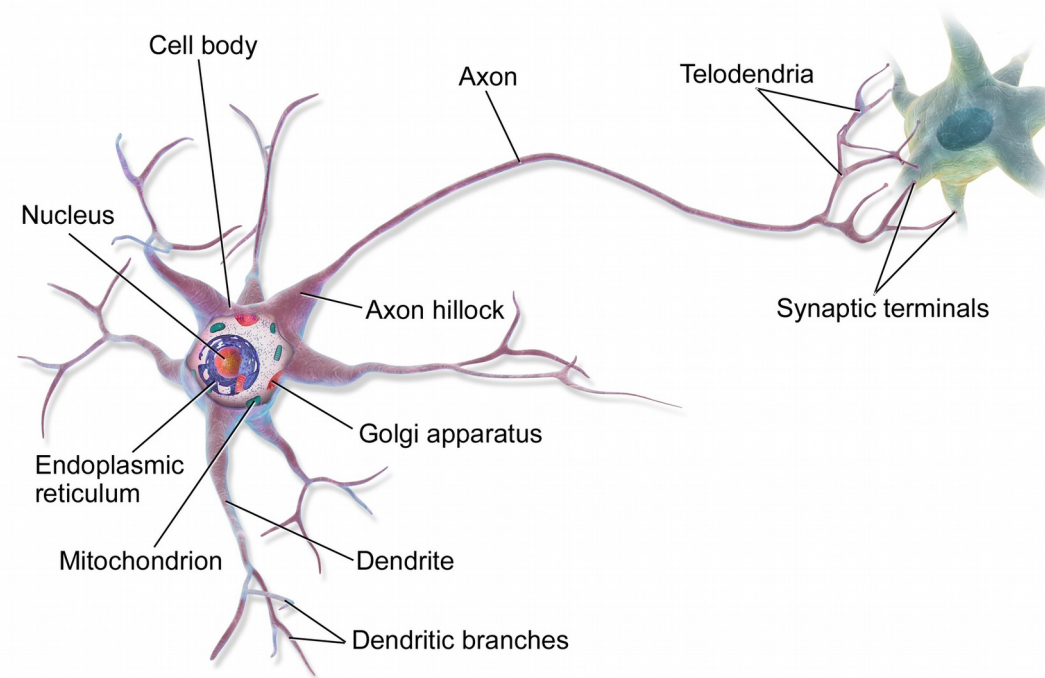
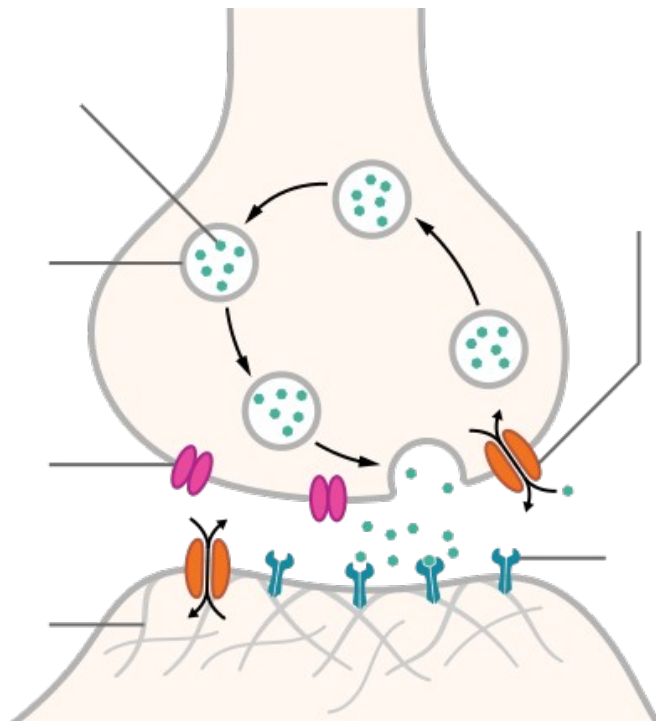
Neurons

- A spike (action potential) needs a threshold influx
- Refractory time limits spike frequency to 200 Hz (~clock rate of the brain)



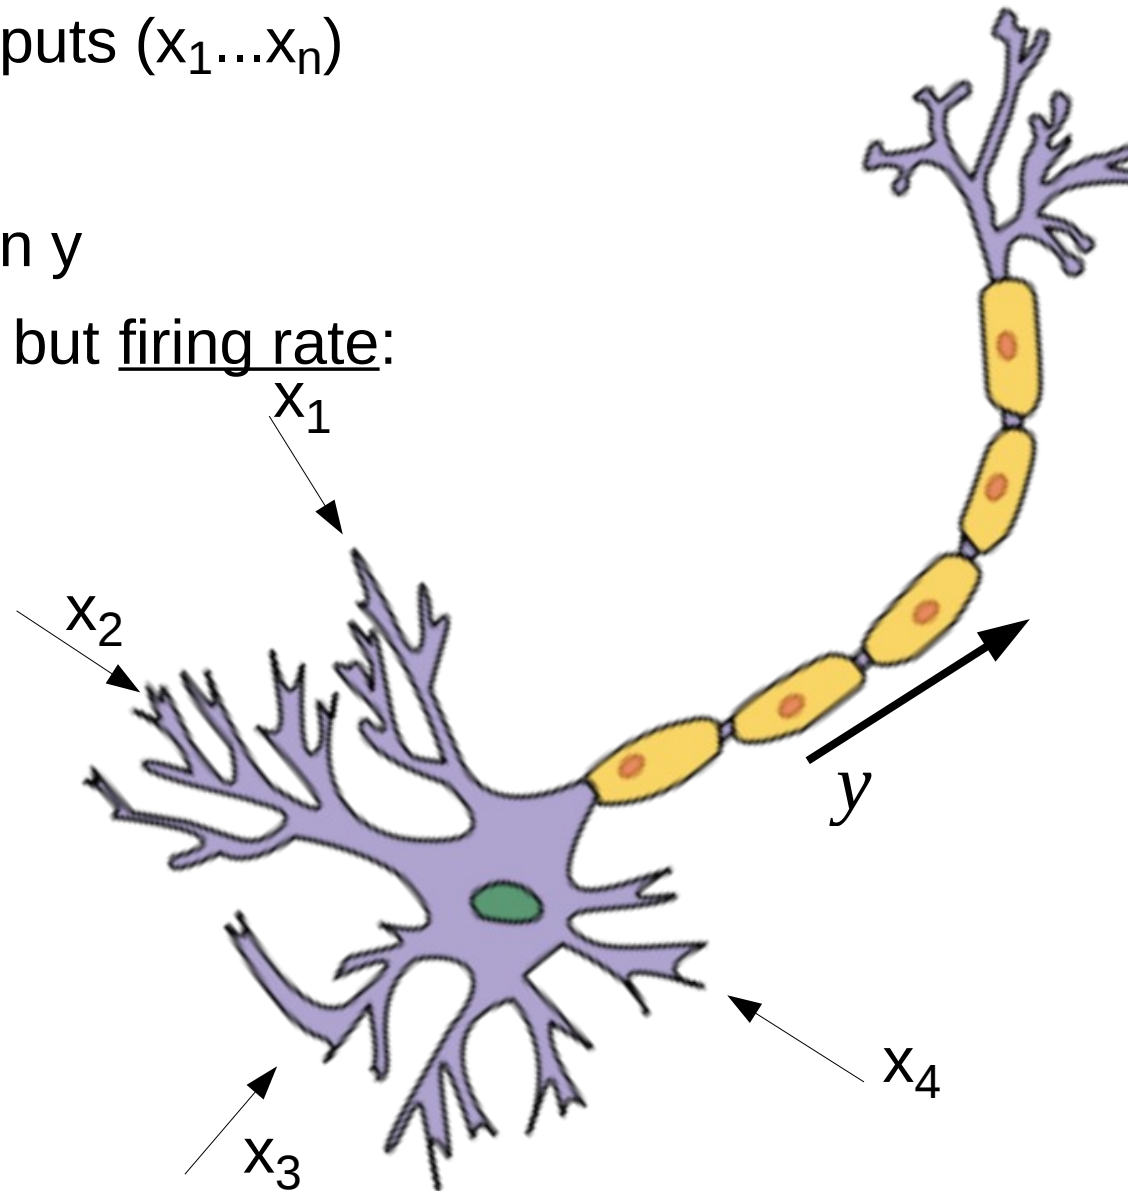
Synapses

- A synapse allows the flow of charge from one neuron to another
- Can “excite” or “inhibit” action potentials
- Synapses change their strength with “experience” (and even grow newly) – “synaptic plasticity”
- Result: changed computation (learning!)



A Simple Model

- A list of input neurons with outputs ($x_1 \dots x_n$)
 - Inputs to this neuron
- Output activation of this neuron y
- Do not code individual spikes, but firing rate:
 - Zero = no action potentials at all
 - One = max. firing frequency



A Simple Model

- Synaptic weights ($w_1 \dots w_n$):

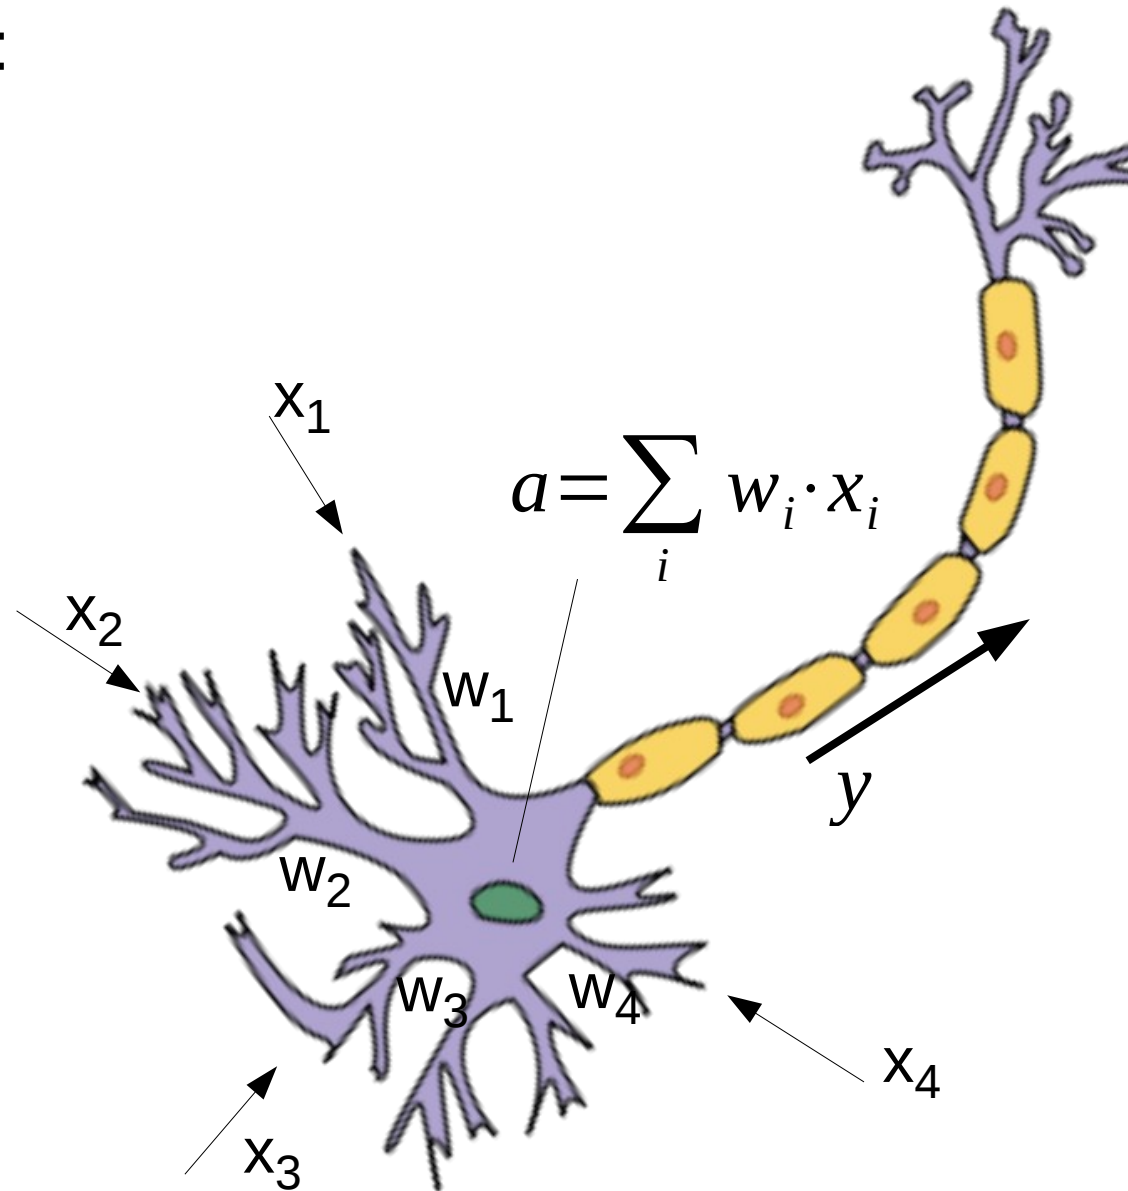
- Positive w = excitatory
- Negative w = inhibitory
- Large = strong effect
- Zero = no effect

- Inflow from neuron i

- Modulated input $w_i \cdot a_i$

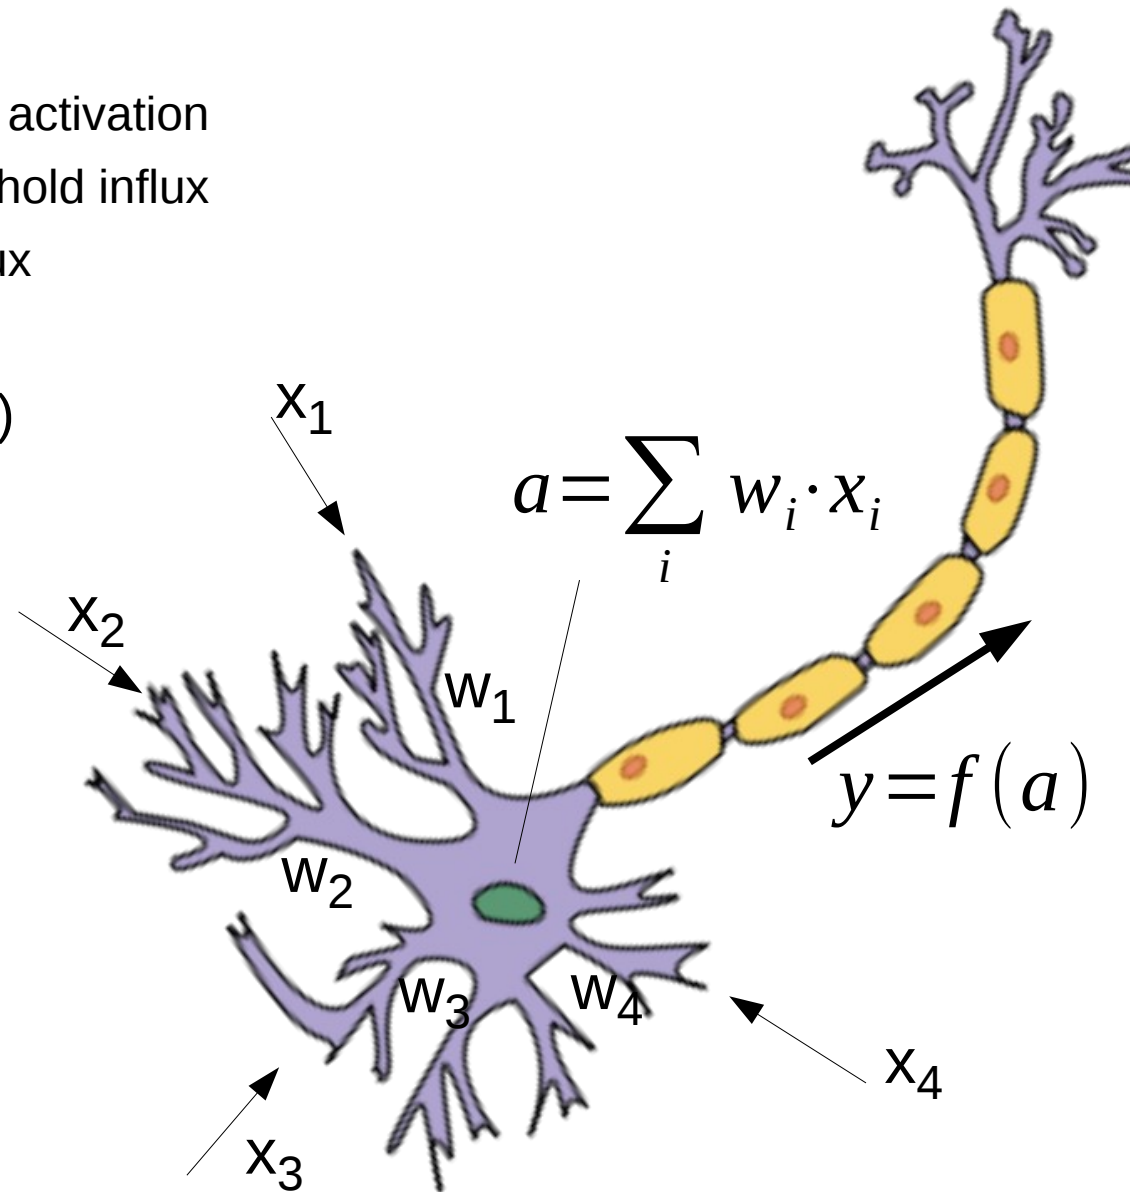
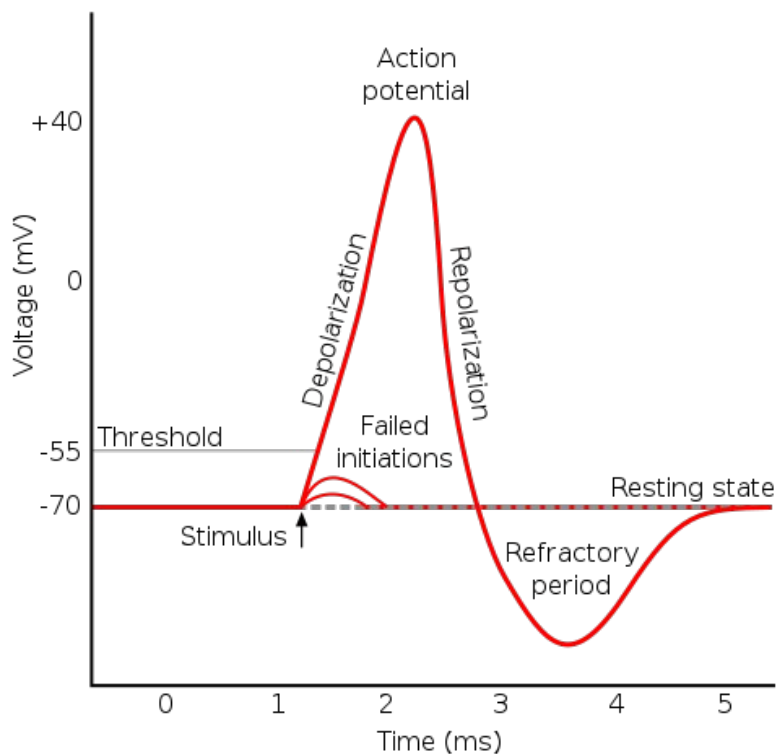
- Net activation a

- Sum of inflows



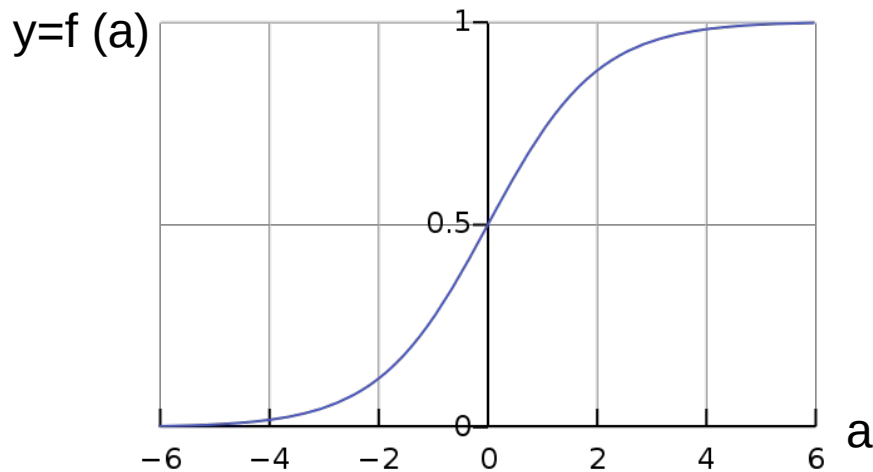
A Simple Model

- Activation function f
 - Determines output firing rate from net activation
 - Low to no output (0.0) for below threshold influx
 - Saturated firing rate (1.0) for high influx
 - Monotonically increasing in between
- Several candidates (design choice)



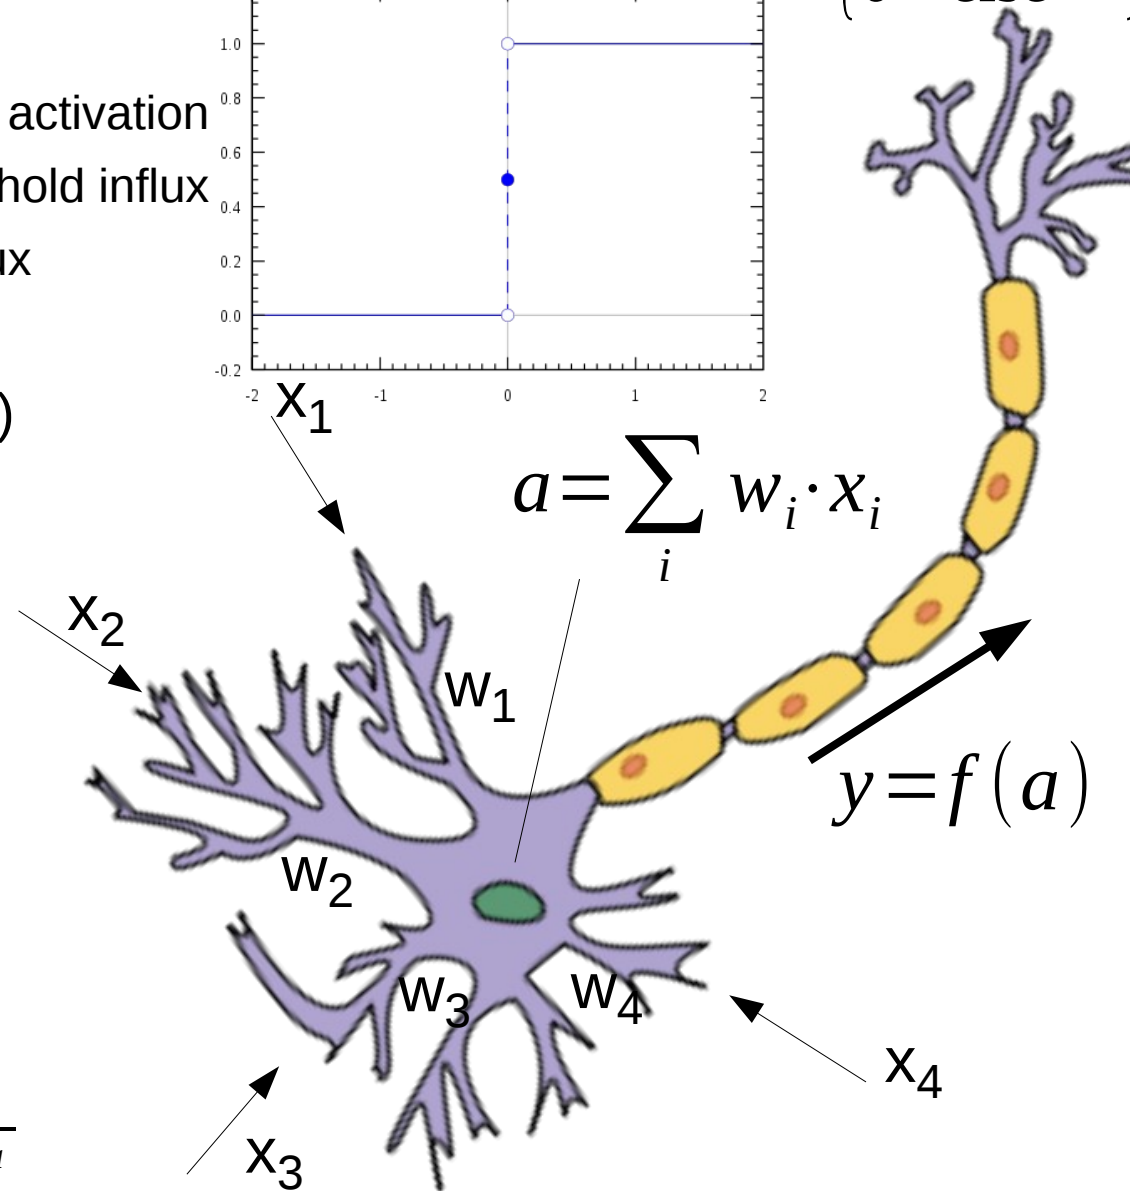
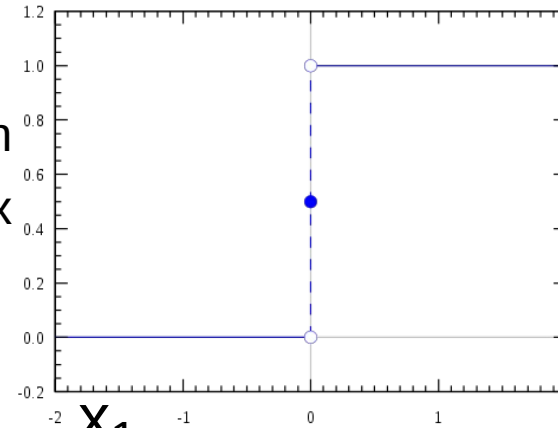
A Simple Model

- Activation function f
 - Determines output firing rate from net activation
 - Low to no output (0.0) for below threshold influx
 - Saturated firing rate (1.0) for high influx
 - Monotonically increasing in between
- Several candidates (design choice)



Sigmoid function:
$$f(a) = \frac{1}{1 + e^{-a}}$$

Step function:
$$f(a) = \begin{cases} 1 & \text{if } a \geq 0 \\ 0 & \text{else} \end{cases}$$



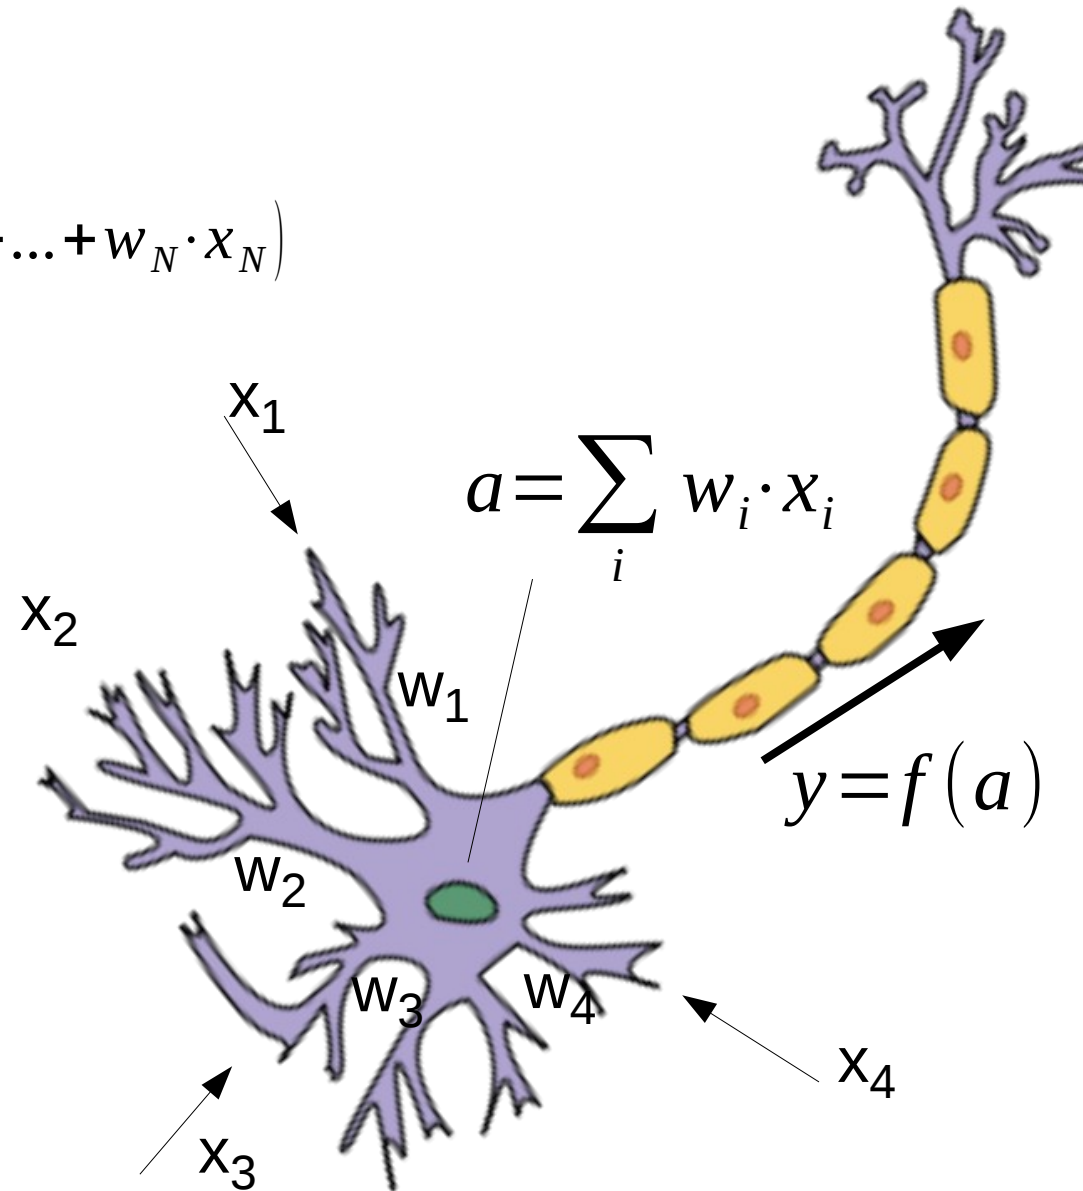
A Simple Model

- In summary:

$$y = f\left(\sum_{i=0}^N w_i \cdot x_i\right) = f(w_0 \cdot x_0 + \dots + w_N \cdot x_N)$$

- In vector notation:

$$y = f\left(\begin{pmatrix} w_0 \\ w_1 \\ \vdots \\ w_2 \end{pmatrix}^T \cdot \begin{pmatrix} 1 \\ x_1 \\ \vdots \\ x_2 \end{pmatrix}\right) = f(\vec{w}^T \cdot \vec{x})$$

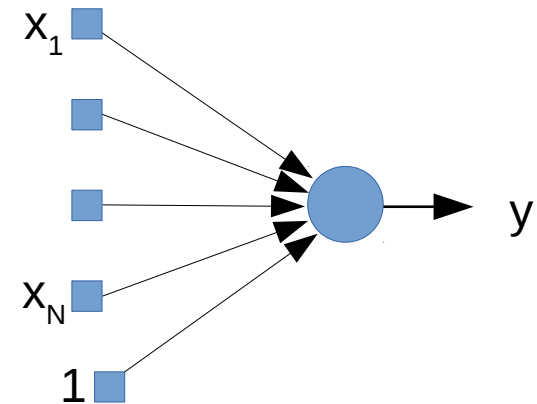


Perceptron

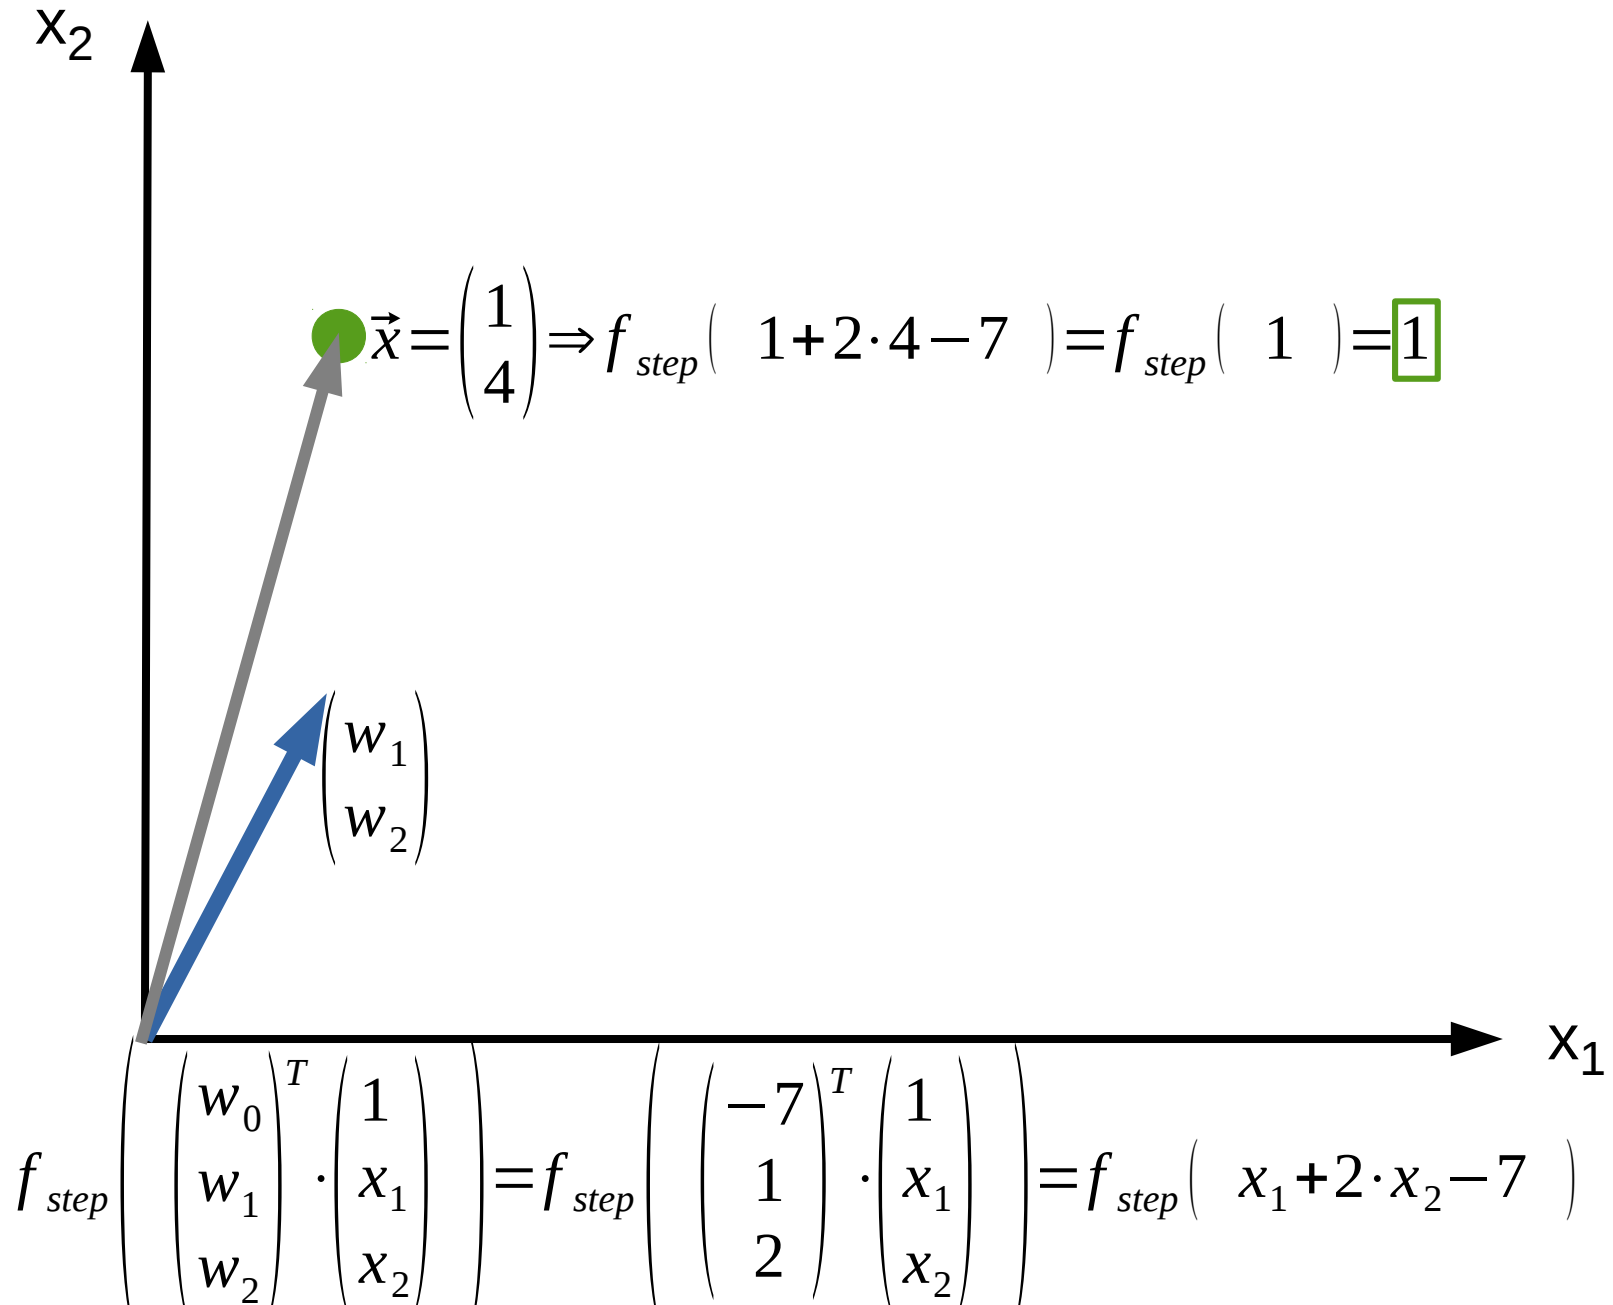
- Invented by Frank Rosenblatt in 1957
- A single neuron with step function
- Supervised learning rule!

$$y = f \left(\begin{pmatrix} w_0 \\ w_1 \\ \vdots \\ w_2 \end{pmatrix}^T \cdot \begin{pmatrix} 1 \\ x_1 \\ \vdots \\ x_2 \end{pmatrix} \right) = f(\vec{w}^T \cdot \vec{x})$$

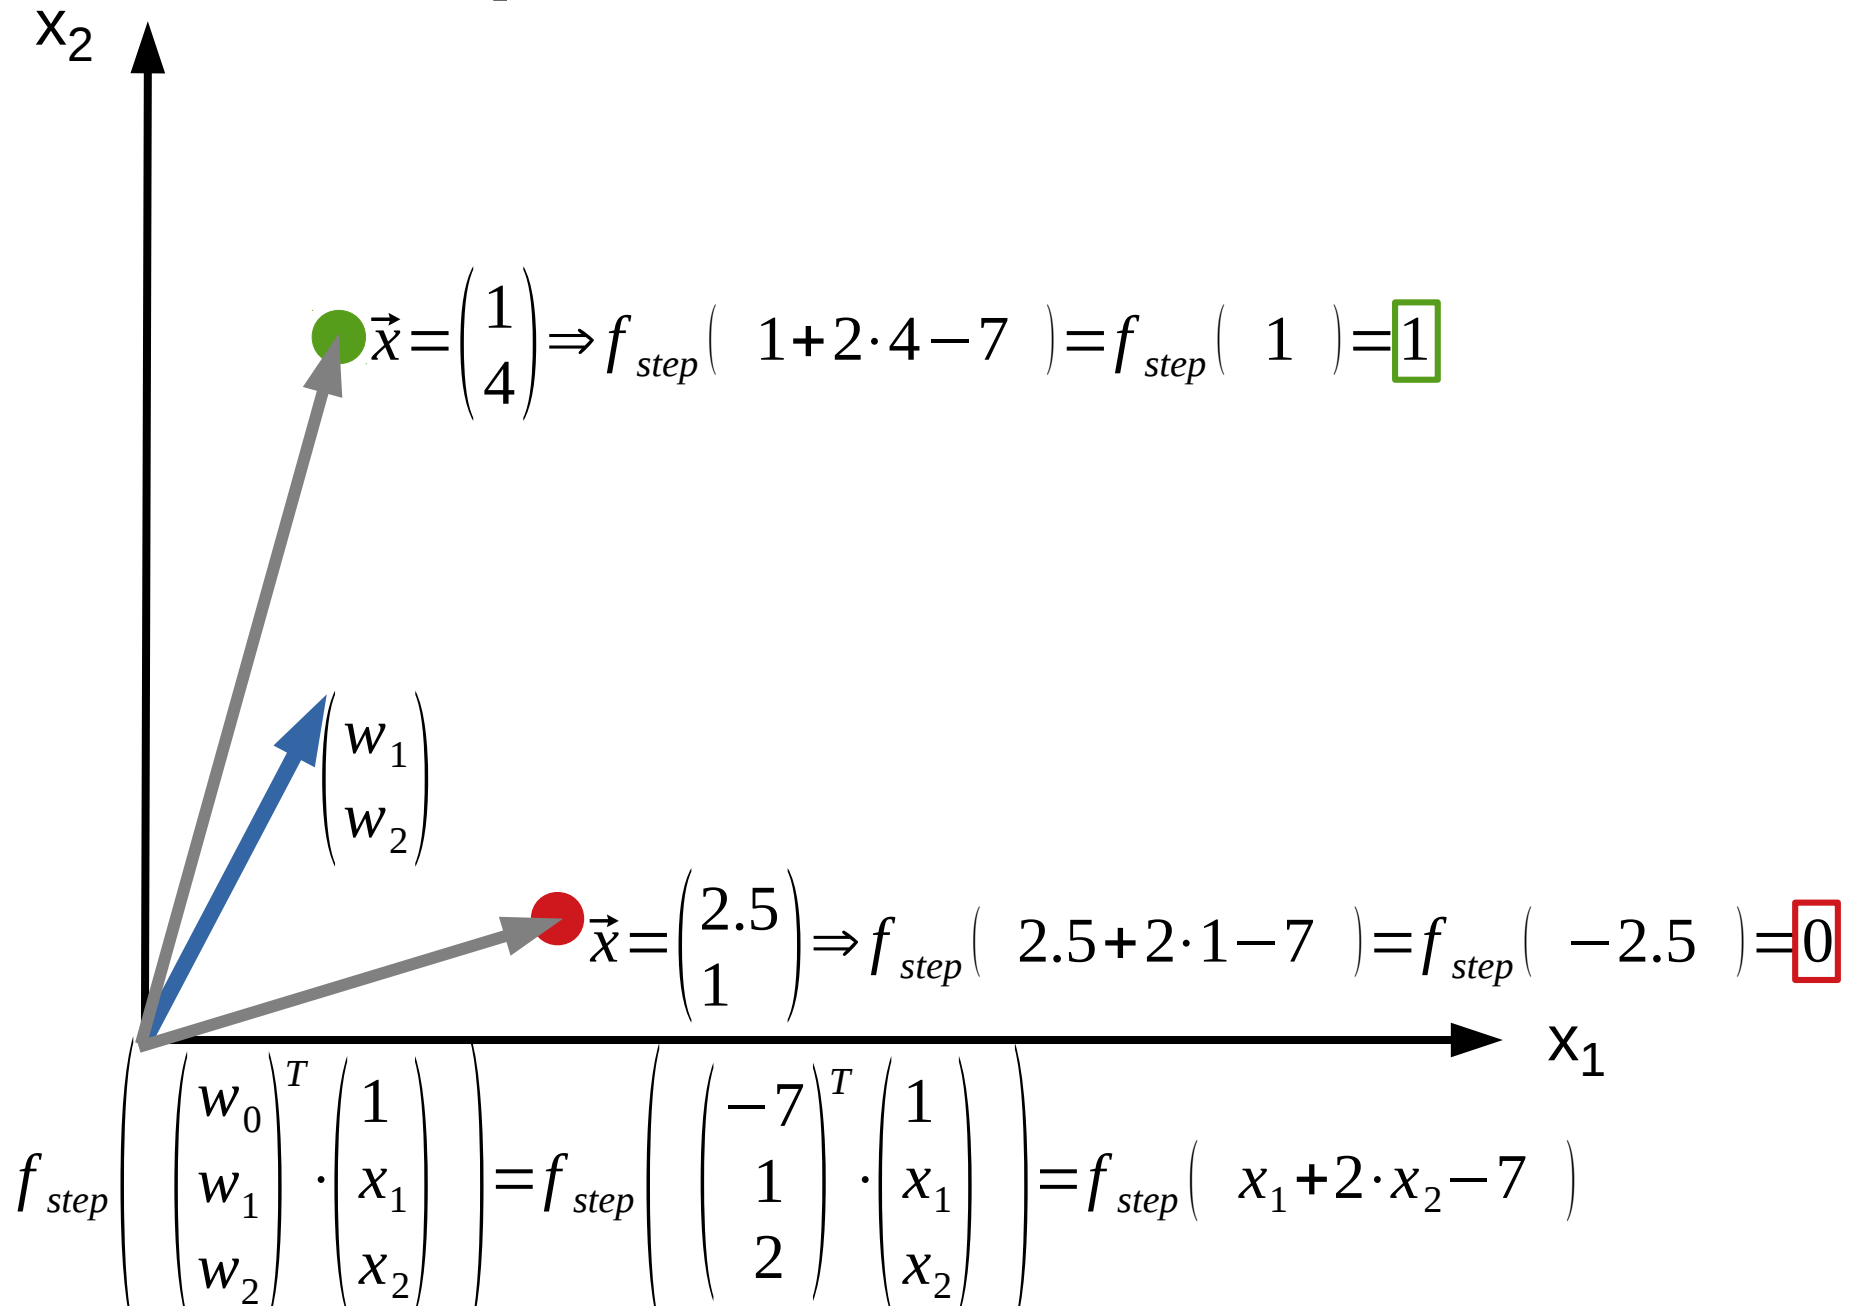
Step function: $f(a) = \begin{cases} 1 & \text{if } a \geq 0 \\ 0 & \text{else} \end{cases}$



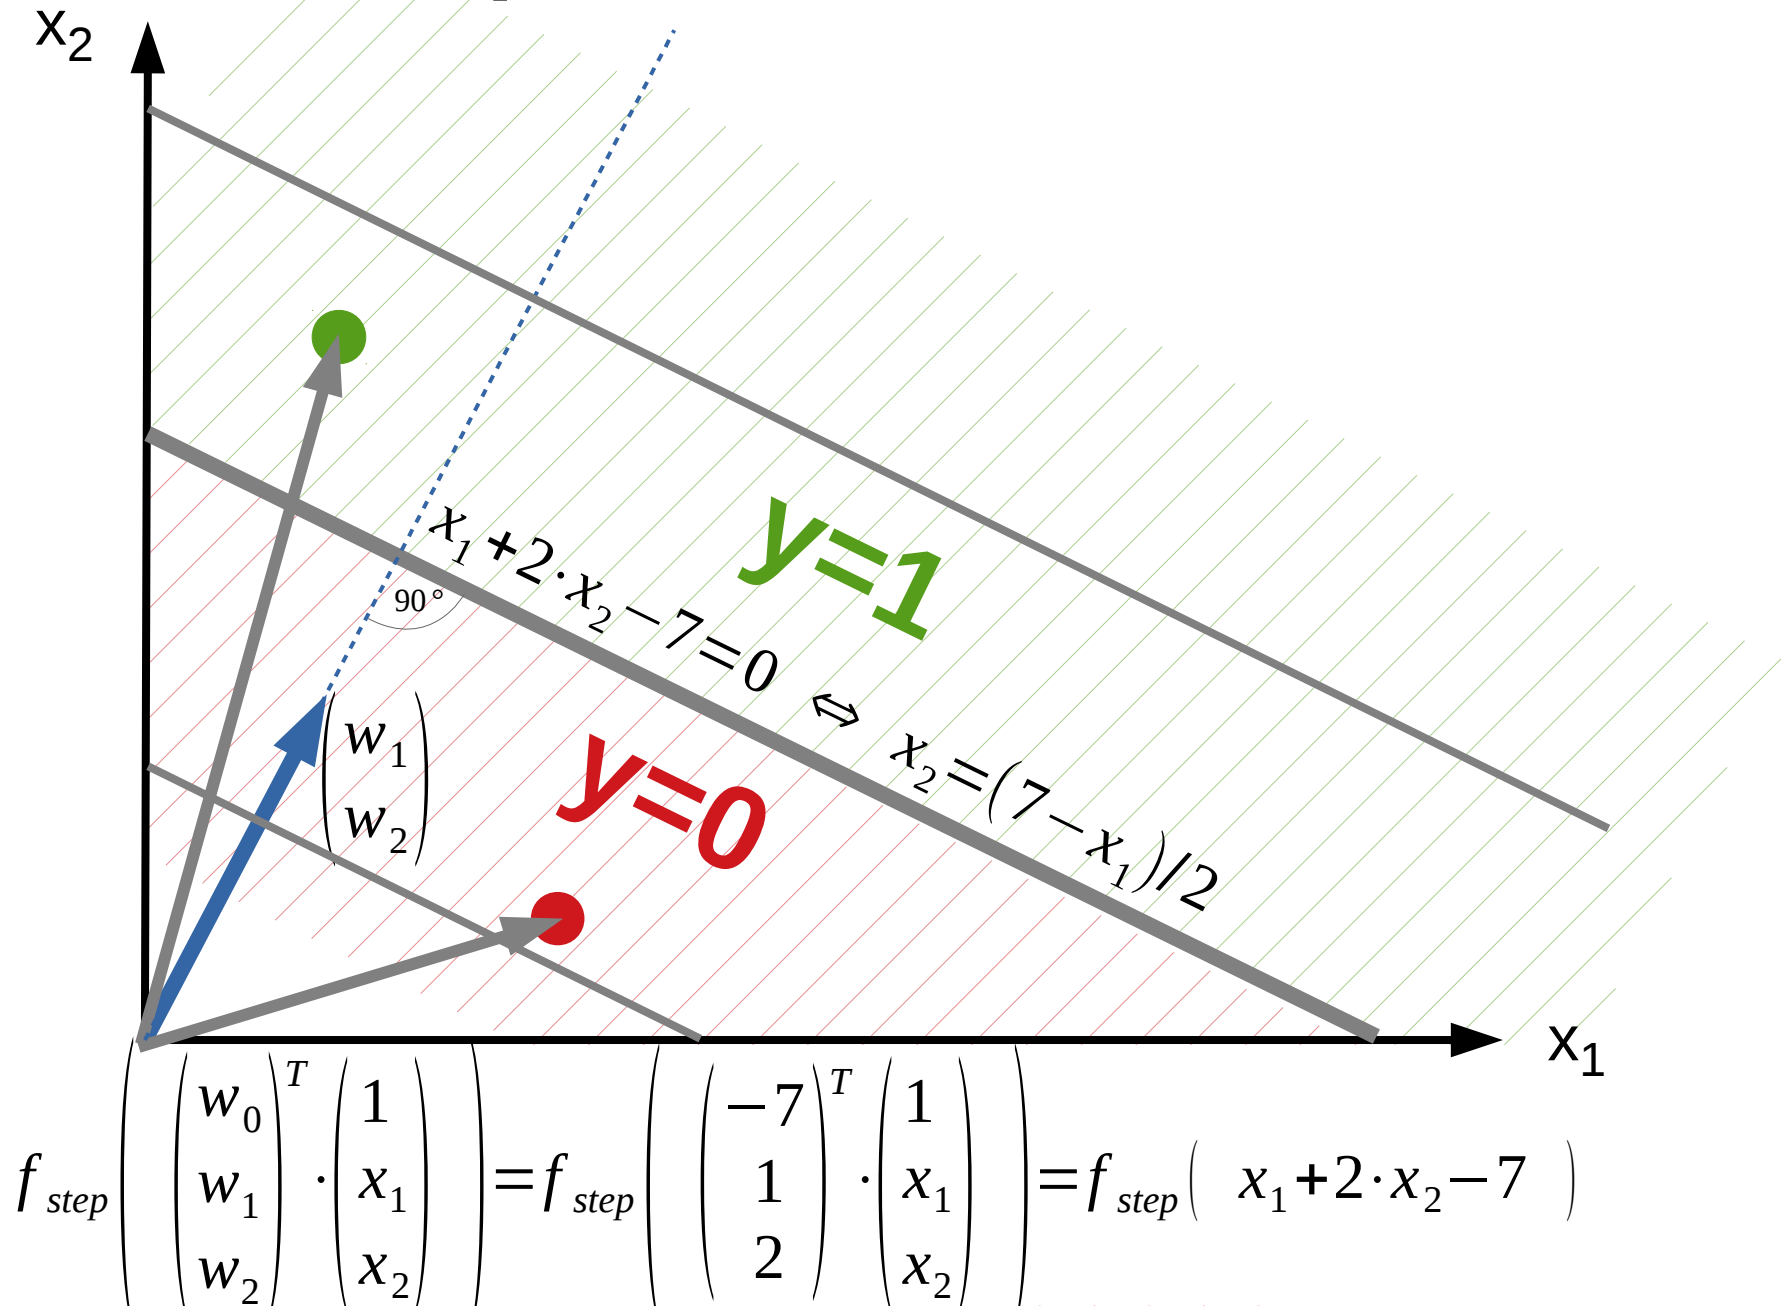
Perceptron Decisions



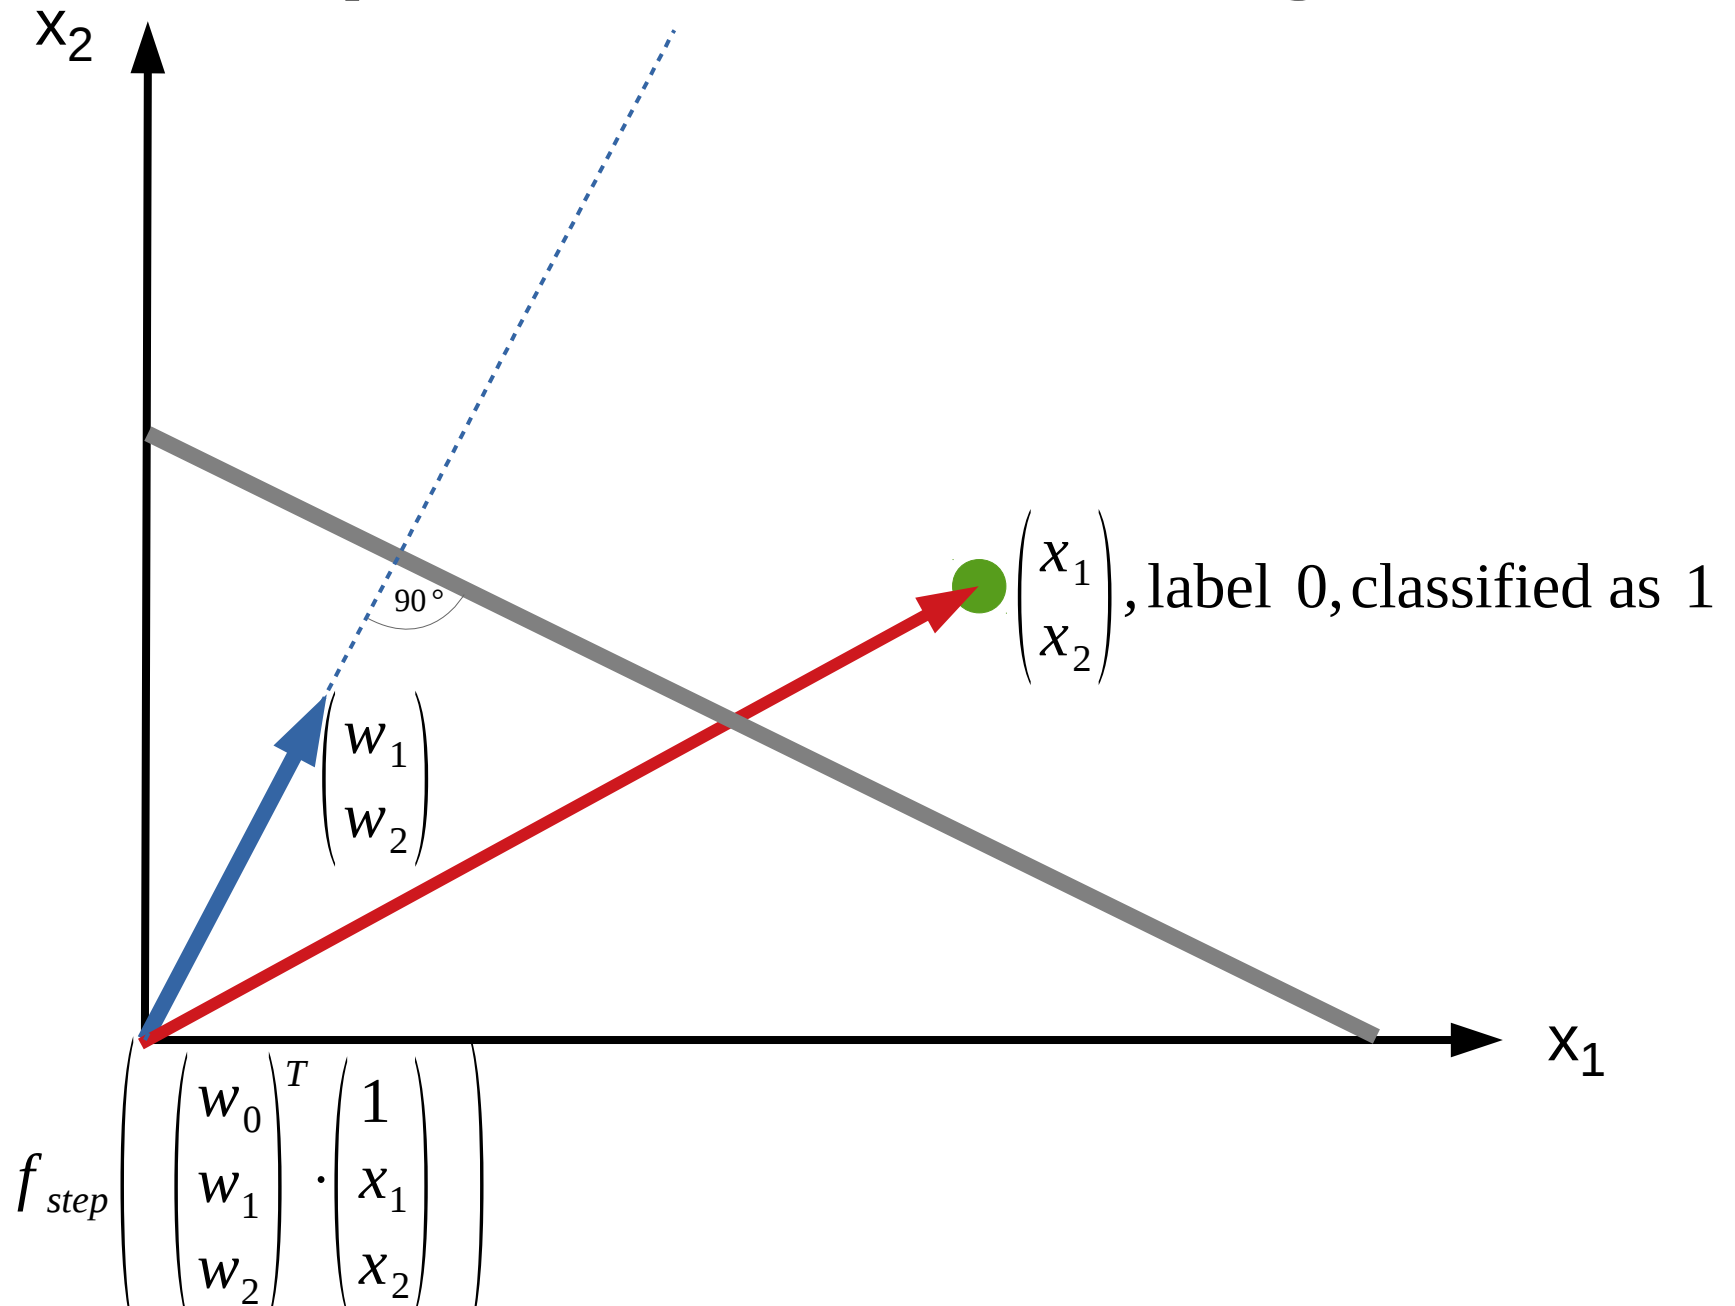
Perceptron Decisions



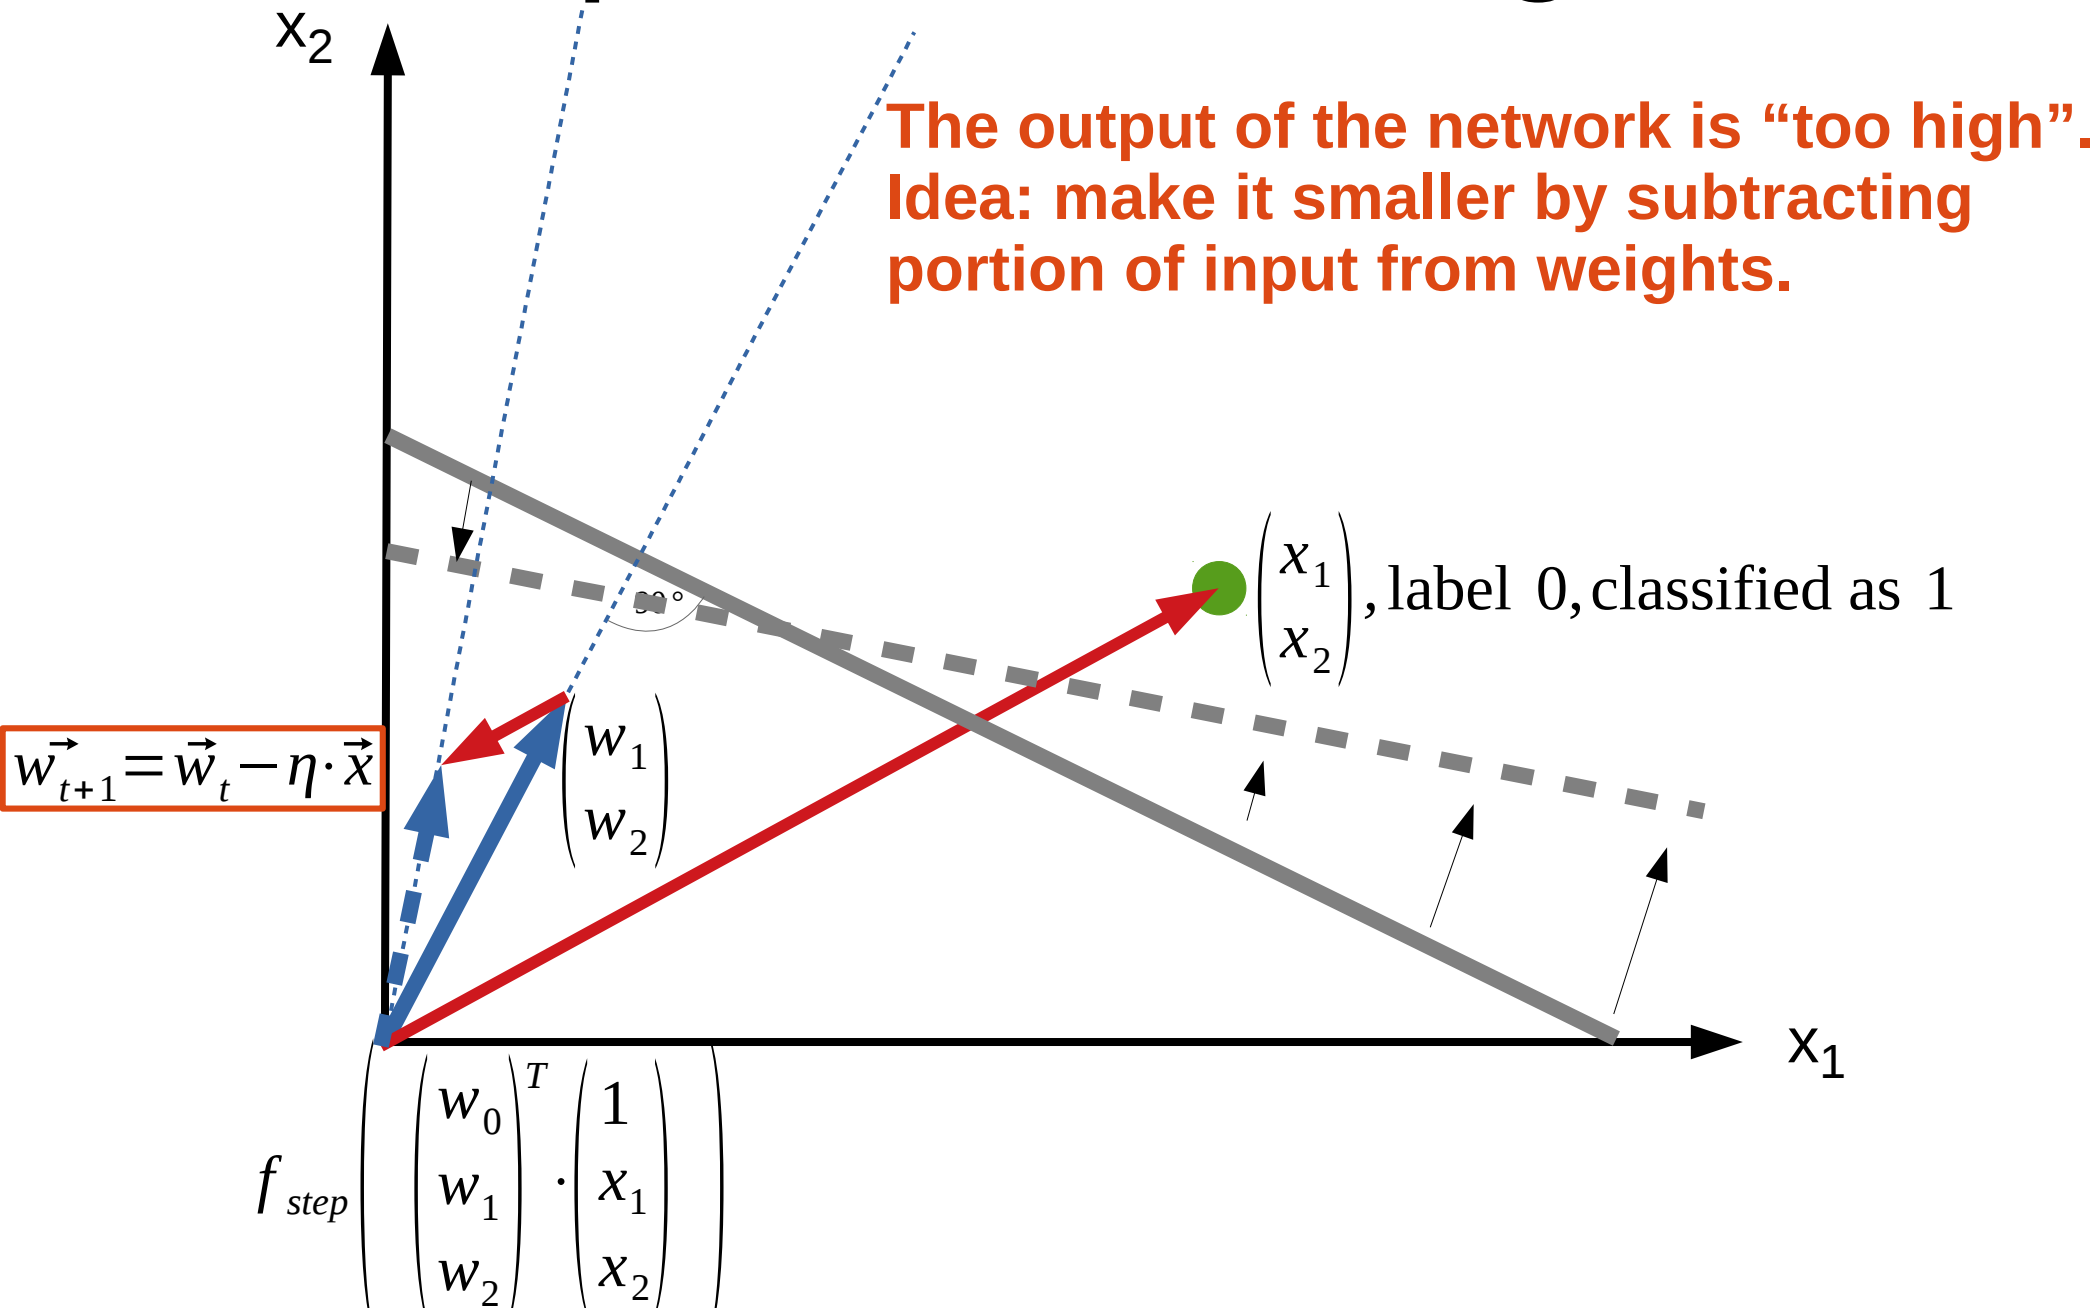
Perceptron Decisions



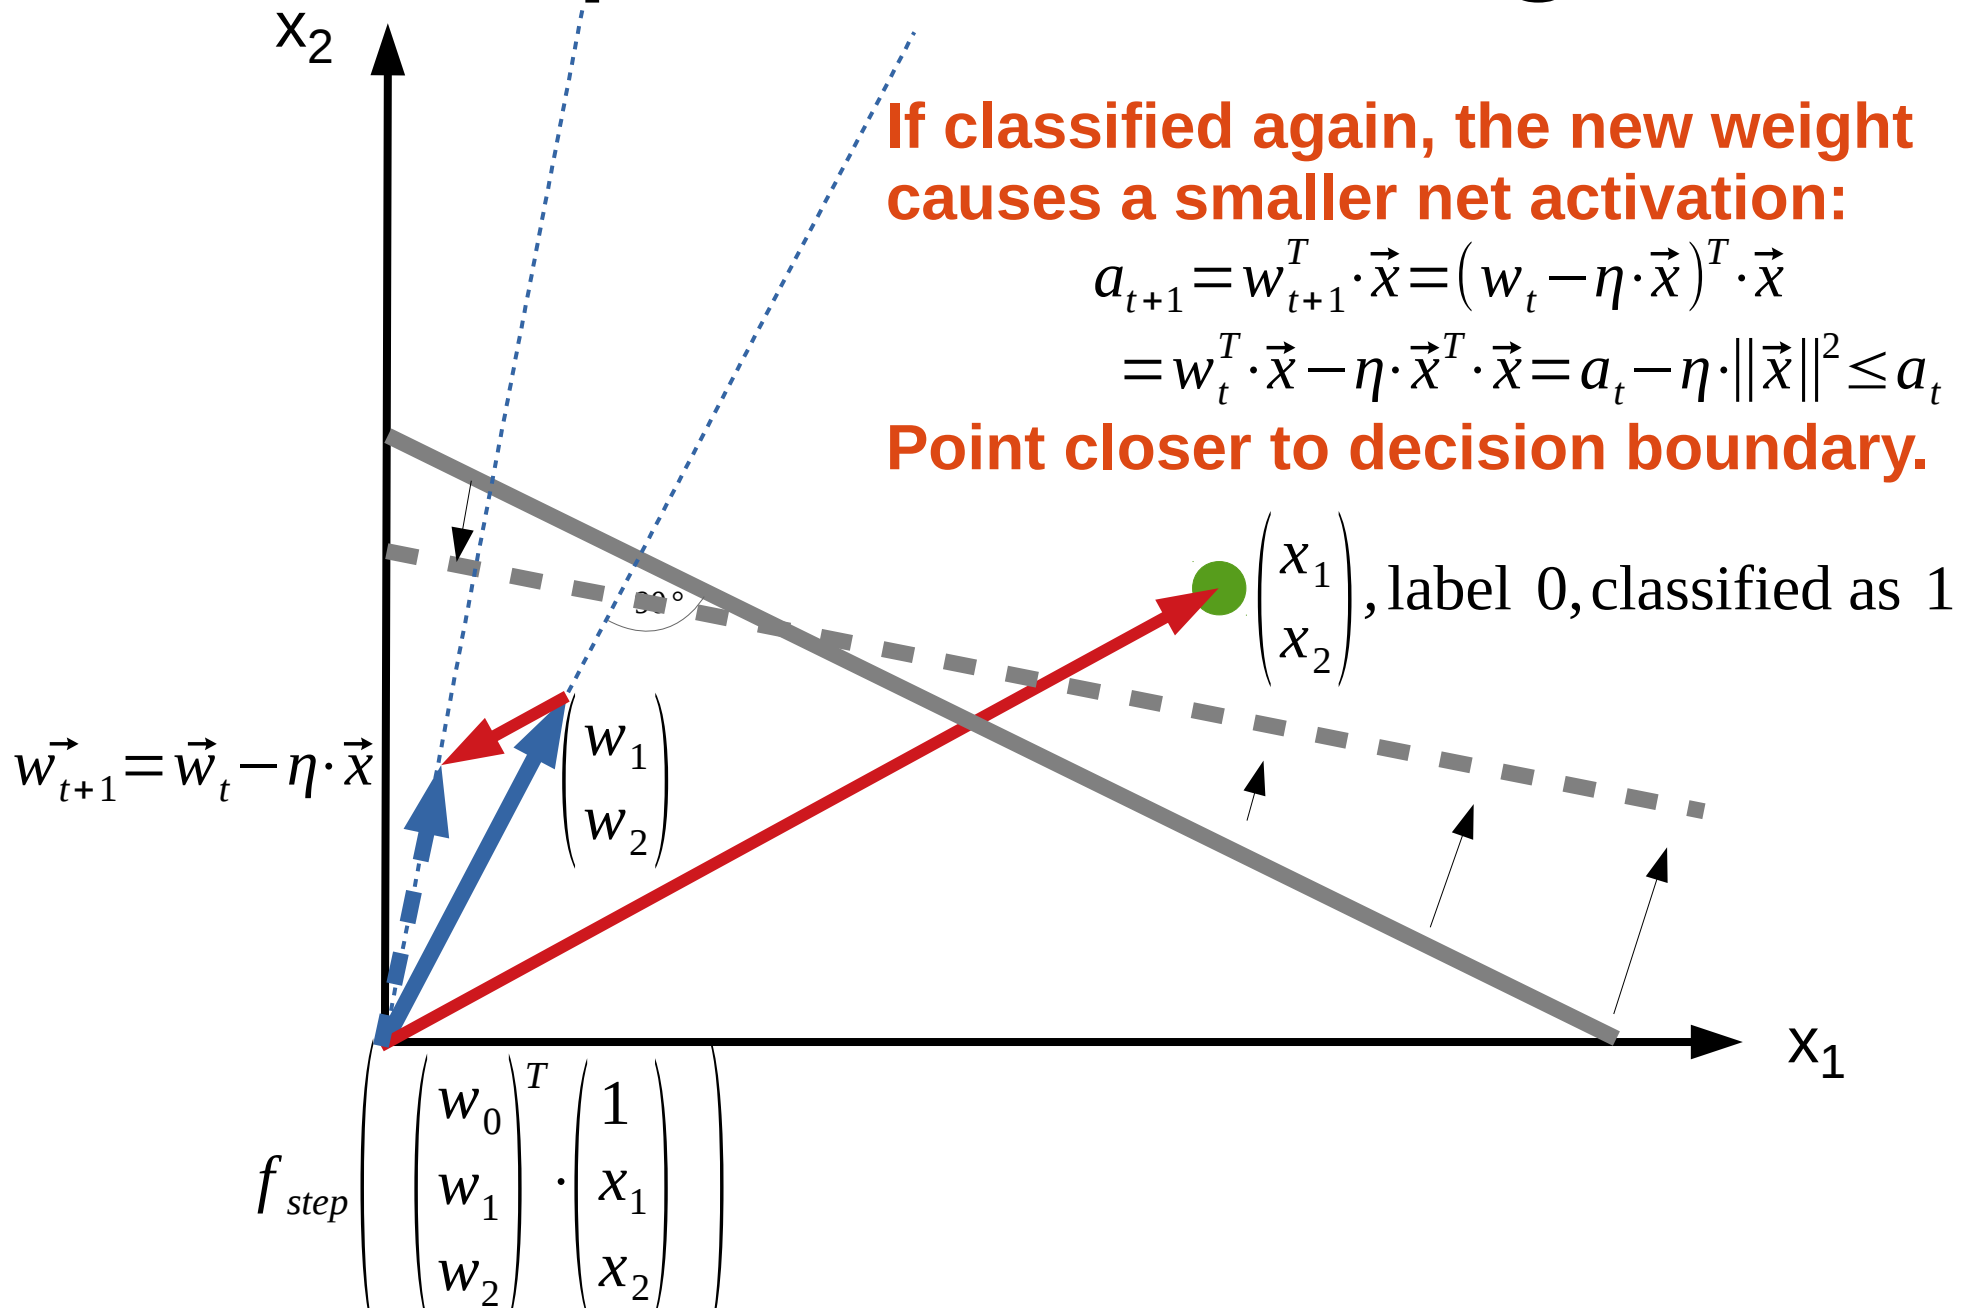
Perceptron Learning Rule



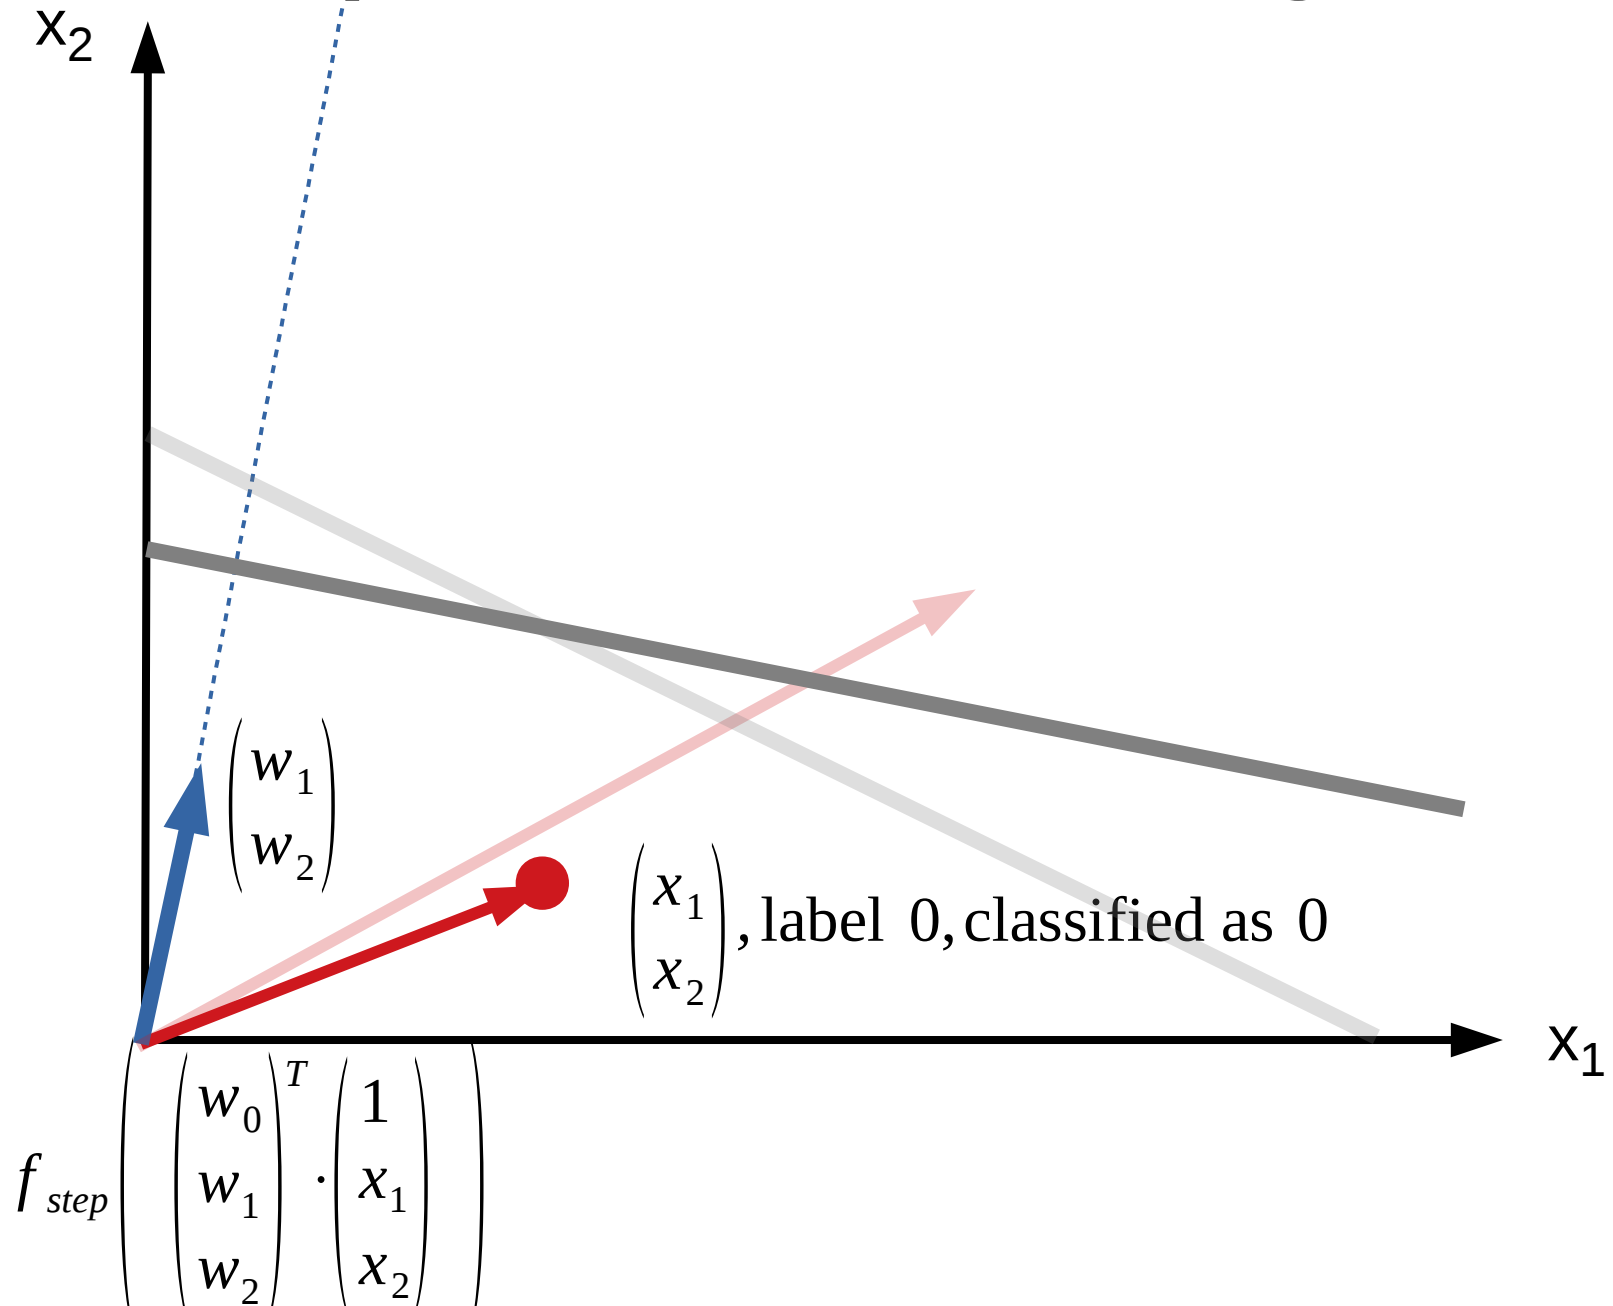
Perceptron Learning Rule



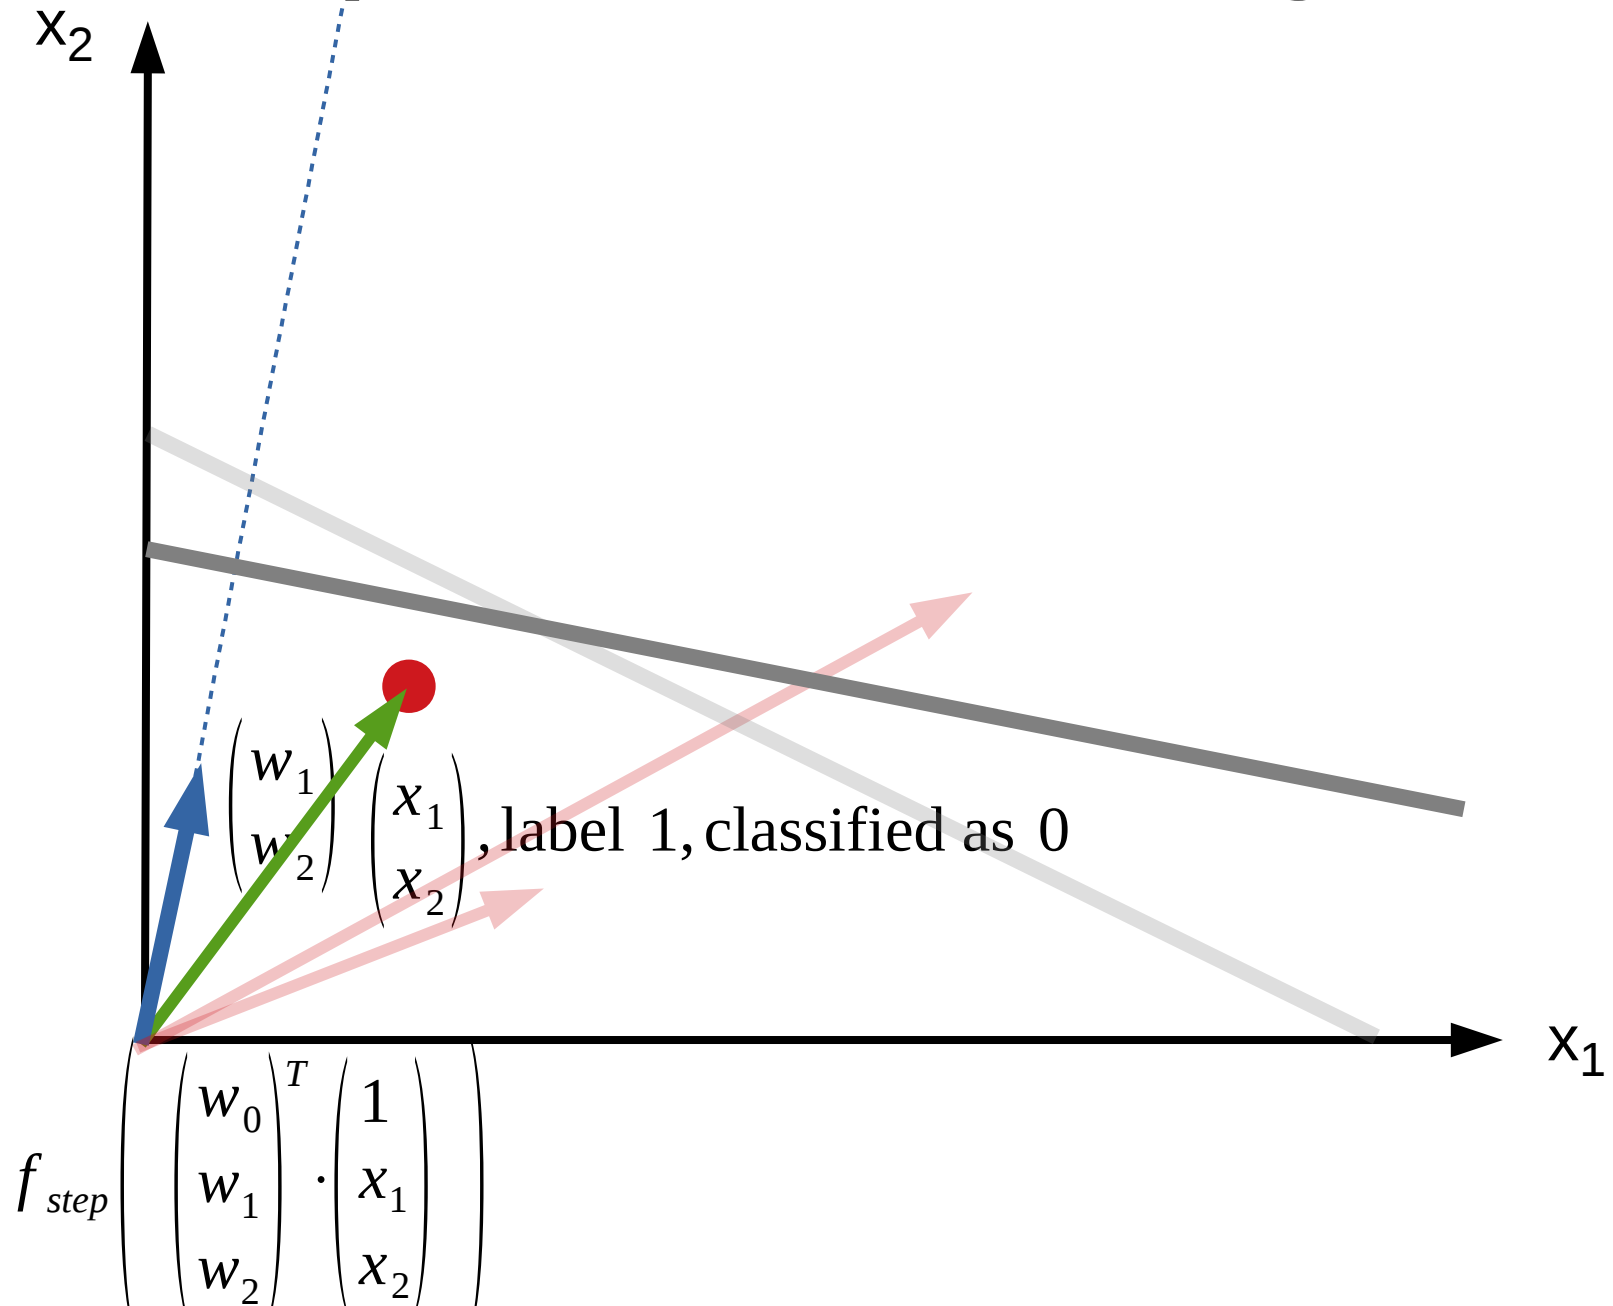
Perceptron Learning Rule



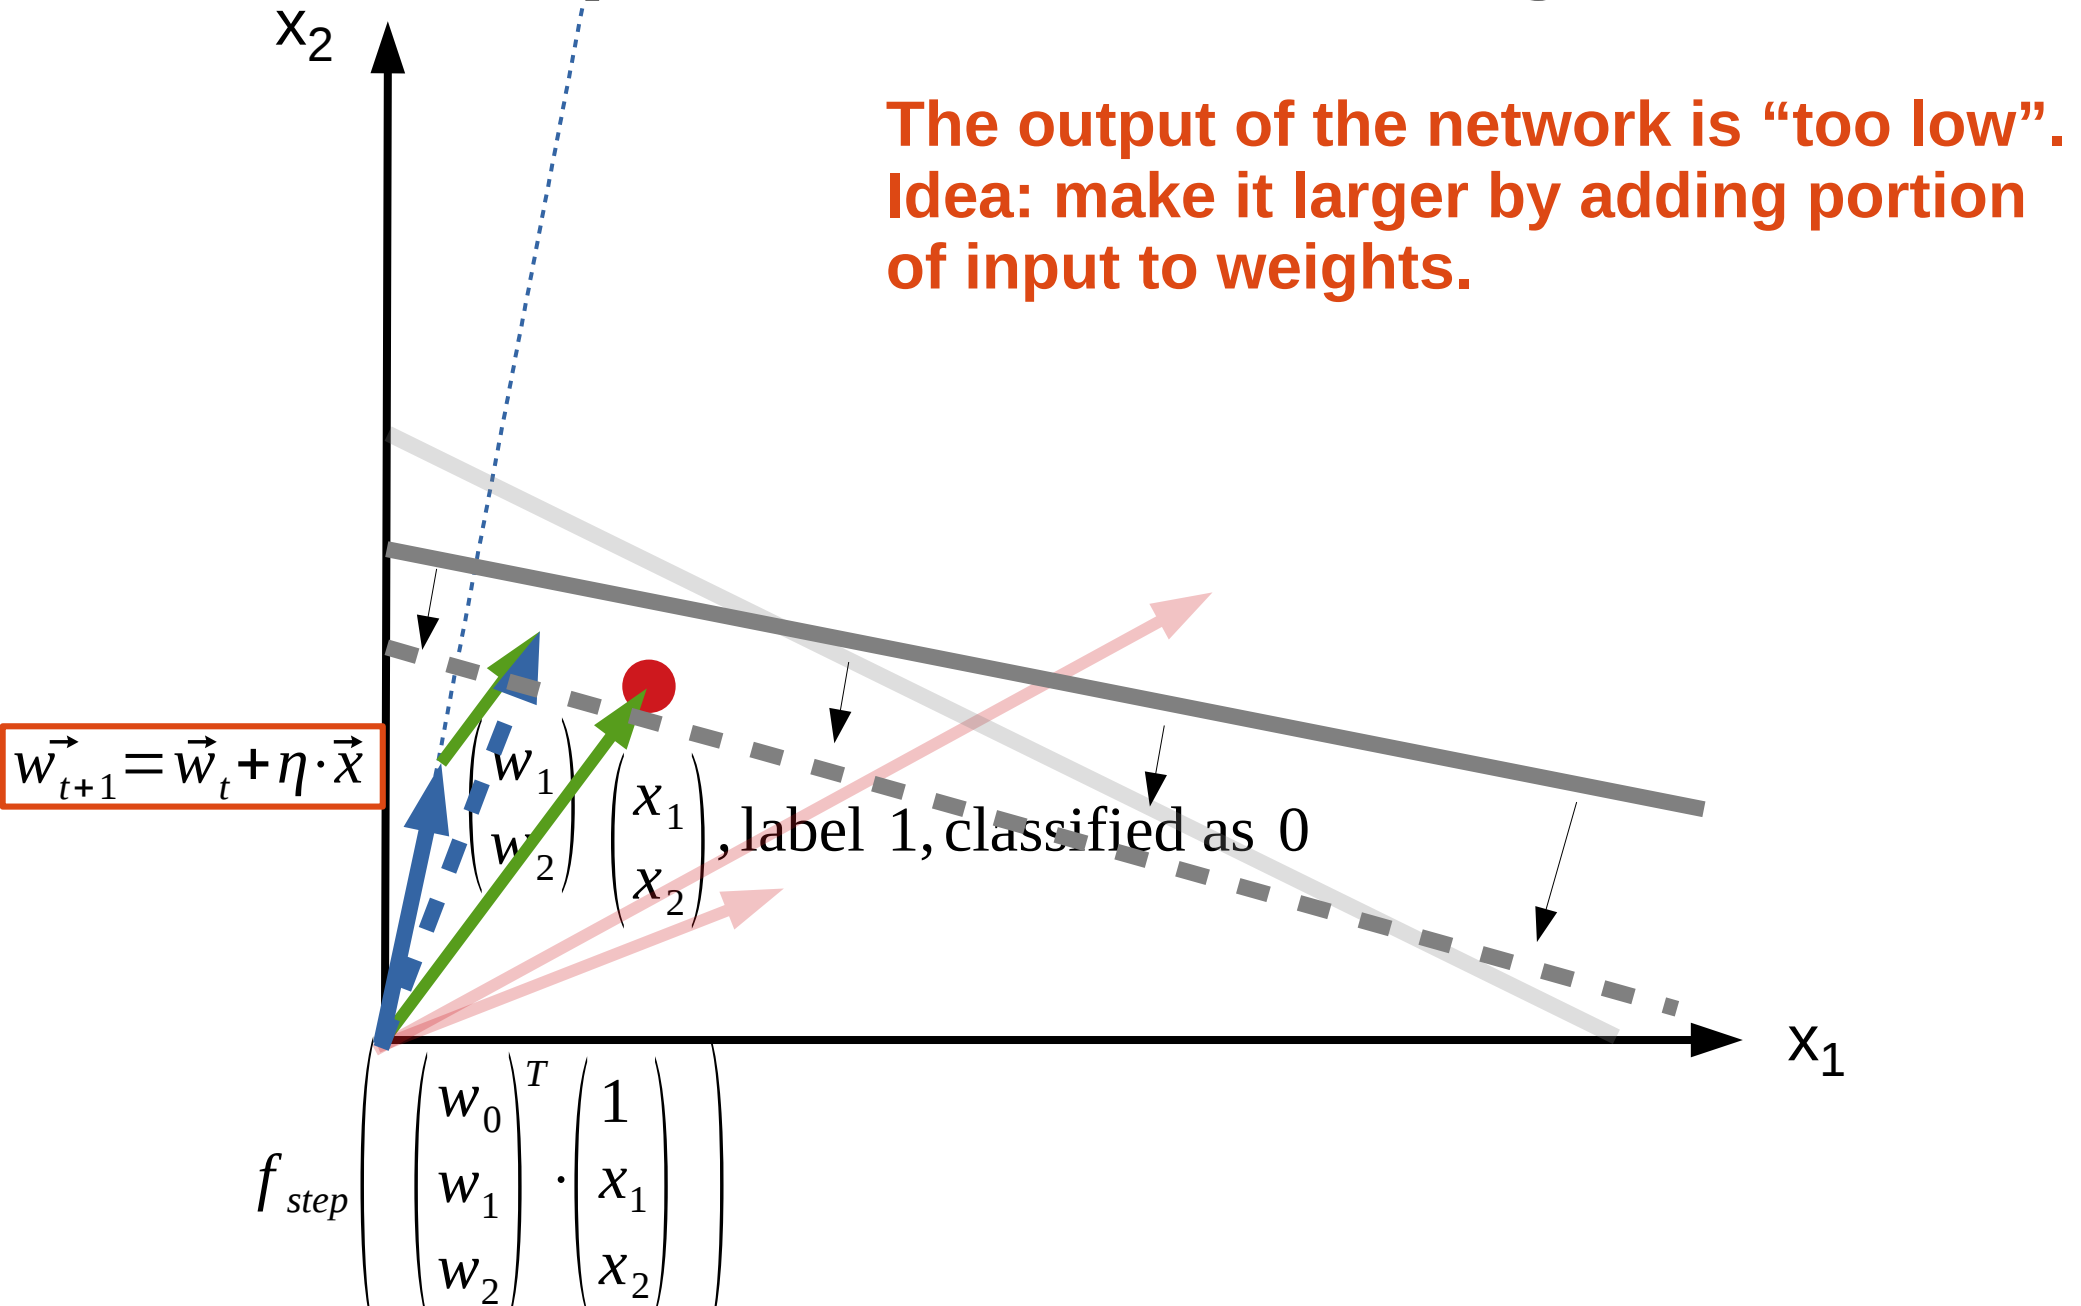
Perceptron Learning Rule



Perceptron Learning Rule



Perceptron Learning Rule



Perceptron Learning Rule

If classified again, the new weight causes a larger net activation:

$$\begin{aligned} a_{t+1} &= w_{t+1}^T \cdot \vec{x} = (w_t + \eta \cdot \vec{x})^T \cdot \vec{x} \\ &= w_t^T \cdot \vec{x} + \eta \cdot \vec{x}^T \cdot \vec{x} = a_t + \eta \cdot \|\vec{x}\|^2 \geq a_t \end{aligned}$$

This example one moves to the correct side right away.

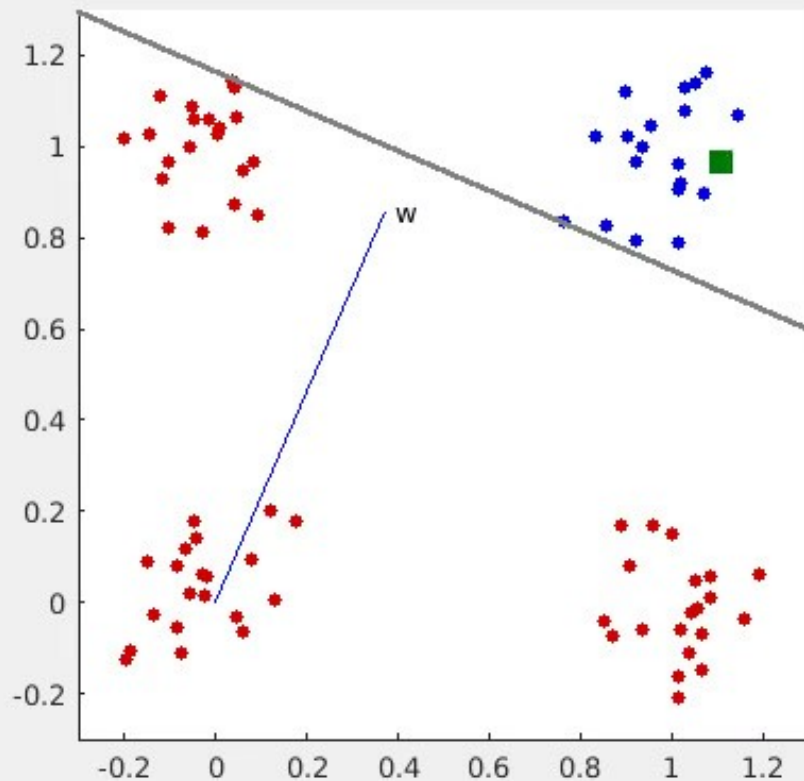
$$\vec{w}_{t+1} = \vec{w}_t + \eta \cdot \vec{x}$$

$\begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$, label 1, classified as 0

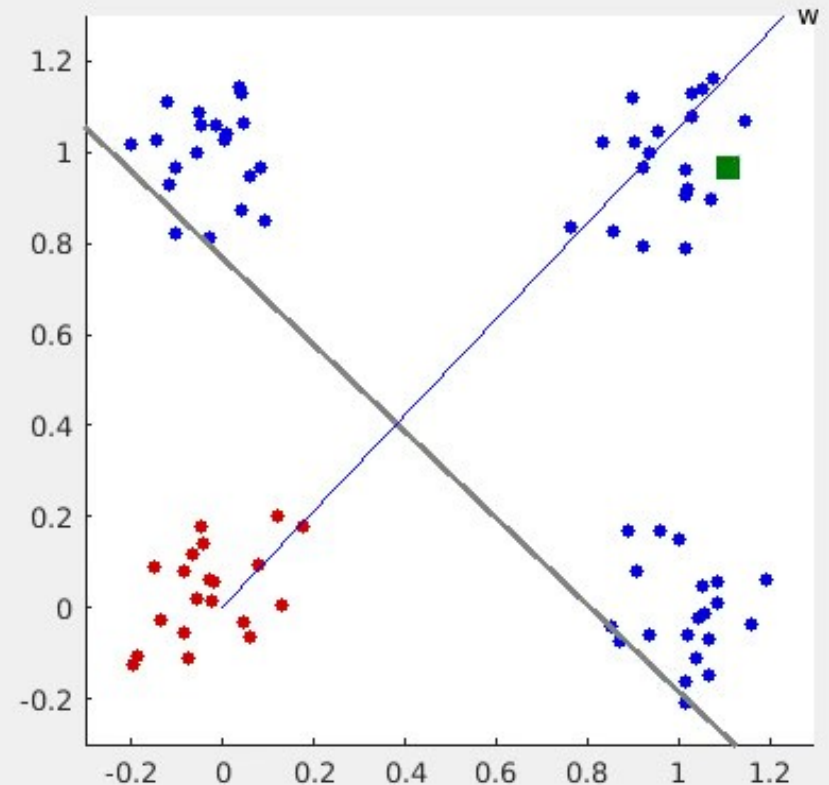
$$f_{step} \begin{pmatrix} w_0^T & 1 \\ w_1 & x_1 \\ w_2 & x_2 \end{pmatrix}$$

Learning Logical Functions

AND



OR



It was by no means previously clear that “neurons” could be “logical”, and even learn to behave so.

This causes a great hype in cognitive science around “connectionism”, where neural networks were used as main metaphor to explain all intelligence

NEW NAVY DEVICE LEARNS BY DOING

Psychologist Shows Embryo
of Computer Designed to
Read and Grow Wiser

WASHINGTON, July 7 (UPI)
—The Navy revealed the embryo of an electronic computer today that it expects will be able to walk, talk, see, write, reproduce itself and be conscious of its existence.

The embryo—the Weather Bureau's \$2,000,000 "704" computer—learned to differentiate between right and left after fifty attempts in the Navy's demonstration for newsmen.

The service said it would use this principle to build the first of its Perceptron thinking machines that will be able to read and write. It is expected to be finished in about a year at a cost of \$100,000.

Dr. Frank Rosenblatt, designer of the Perceptron, conducted the demonstration. He said the machine would be the first device to think as the human brain. As do human be-

ings, Perceptron will make mistakes at first, but will grow wiser as it gains experience, he said.

Dr. Rosenblatt, a research psychologist at the Cornell Aeronautical Laboratory, Buffalo, said Perceptrons might be fired to the planets as mechanical space explorers.

Without Human Controls

The Navy said the perceptron would be the first non-living mechanism "capable of receiving, recognizing and identifying its surroundings without any human training or control."

The "brain" is designed to remember images and information it has perceived itself. Ordinary computers remember only what is fed into them on punch cards or magnetic tape.

Later Perceptrons will be able to recognize people and call out their names and instantly translate speech in one language to speech or writing in another language, it was predicted.

Mr. Rosenblatt said in principle it would be possible to build brains that could reproduce themselves on an assembly line and which would be conscious of their existence.

1958 New York
Times...

In today's demonstration, the "704" was fed two cards, one with squares marked on the left side and the other with squares on the right side.

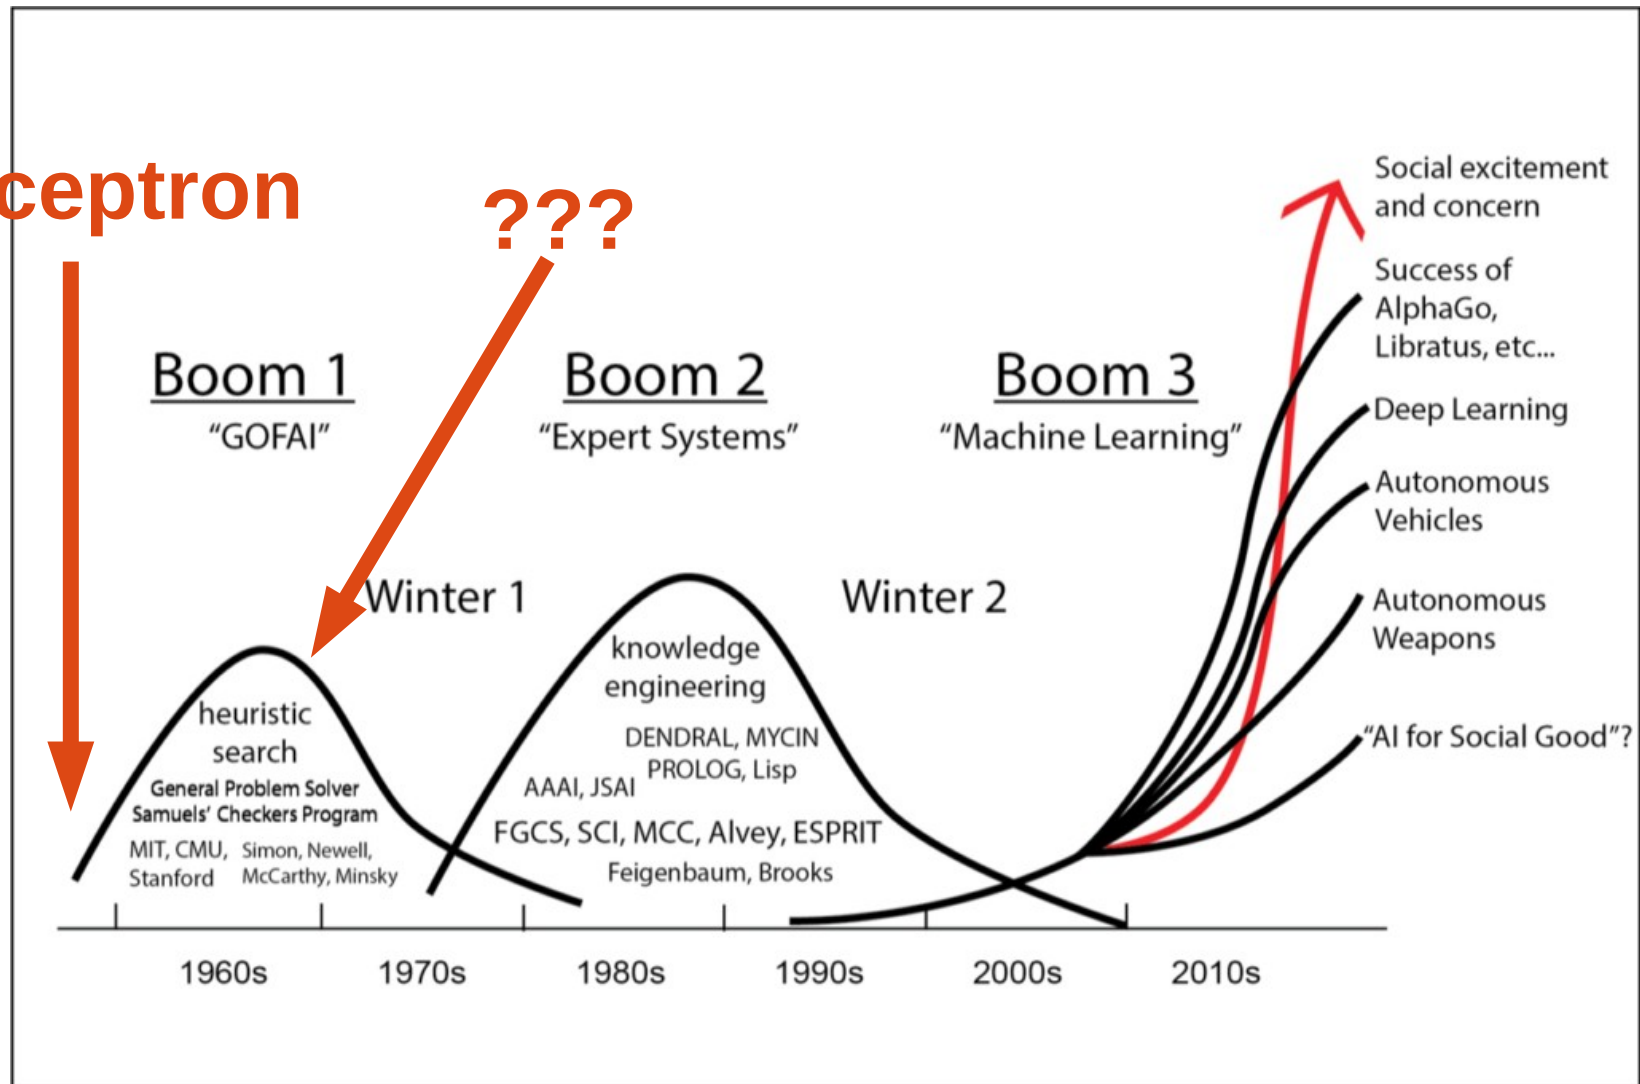
Learns by Doing

In the first fifty trials, the machine made no distinction between them. It then started registering a "Q" for the left squares and "O" for the right squares.

Dr. Rosenblatt said he could explain why the machine learned only in highly technical terms. But he said the computer had undergone a "self-induced change in the wiring diagram."

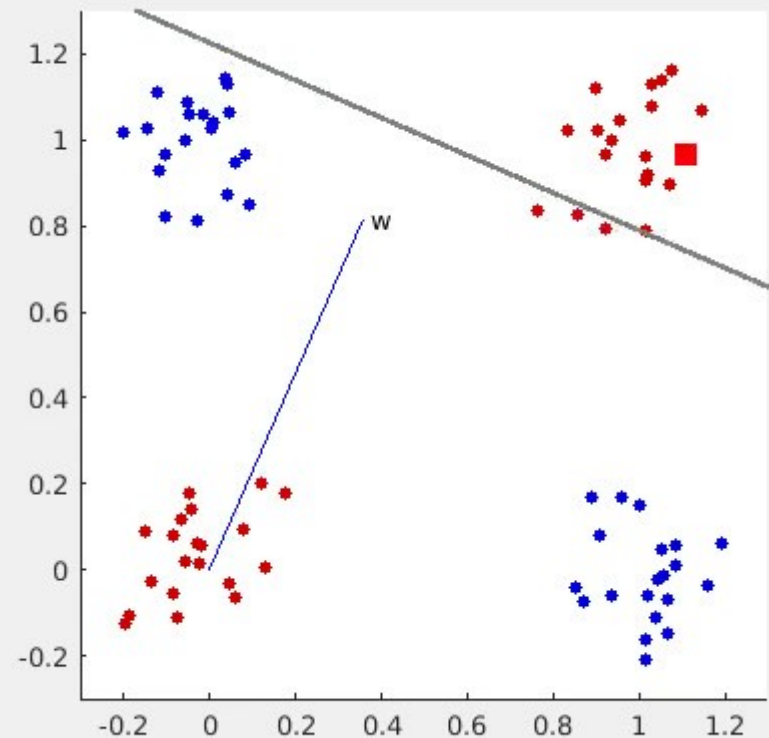
The first Perceptron will have about 1,000 electronic "association cells" receiving electrical impulses from an eye-like scanning device with 400 photo-cells. The human brain has 10,000,000,000 responsive cells, including 100,000,000 connections with the eyes.

Perceptron



The XOR Shock

- In 1969 Marvin Minsky and Seymour Papert (famous AI pioneers) published this example
- XOR = Exclusive “OR”

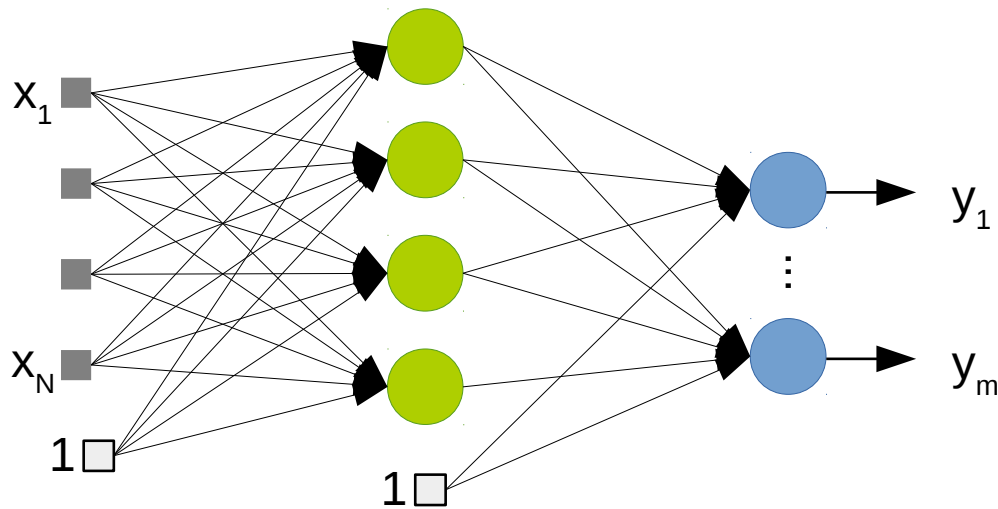


The XOR Shock

- Single neurons are fundamentally unable to describe an XOR function
- Its classes are not linearly separable
 - No possible straight line could separate them
- We need something non-linear...
- Features, maybe?
- Neural features?

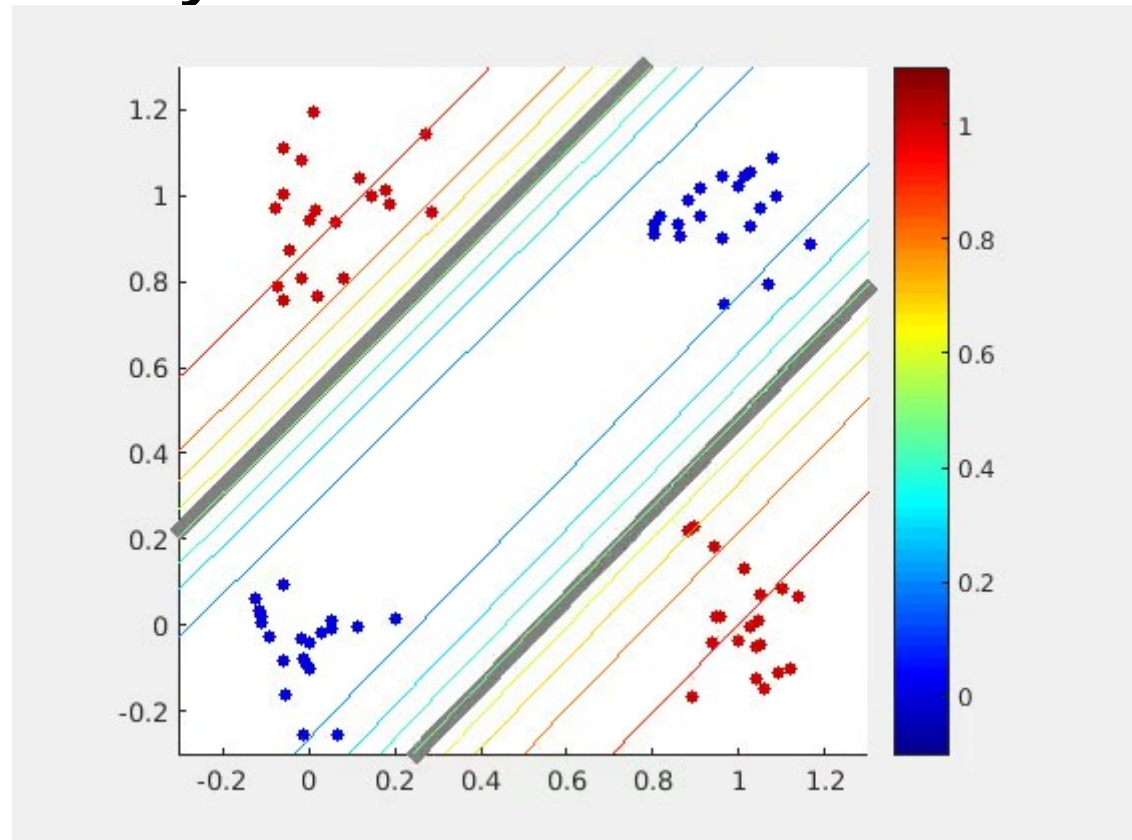
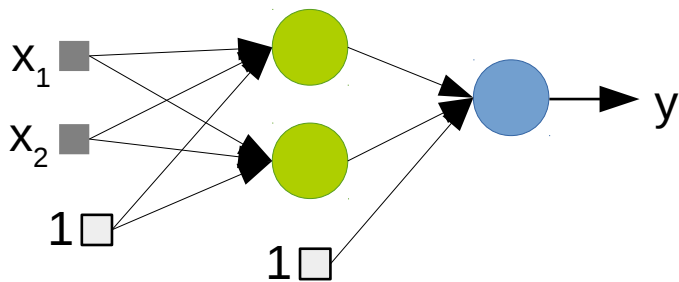
Multi-layer perceptron

- Modern neural networks are layered
- Layers of neurons provide non-linear features for the next ones

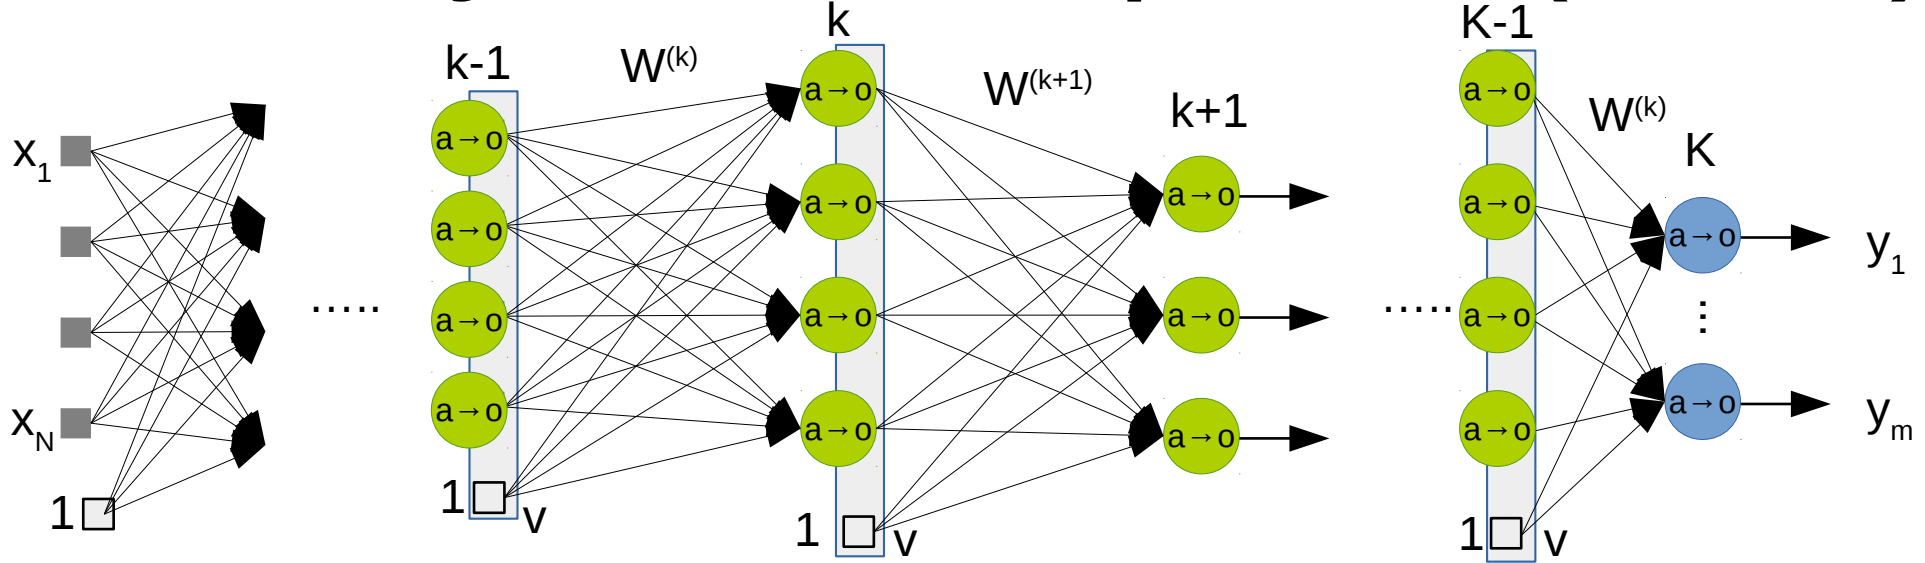


Multi-layer perceptron

- A simple two layer solution to XOR:
- One hidden neuron each for one red cluster
- “or” in the output layer

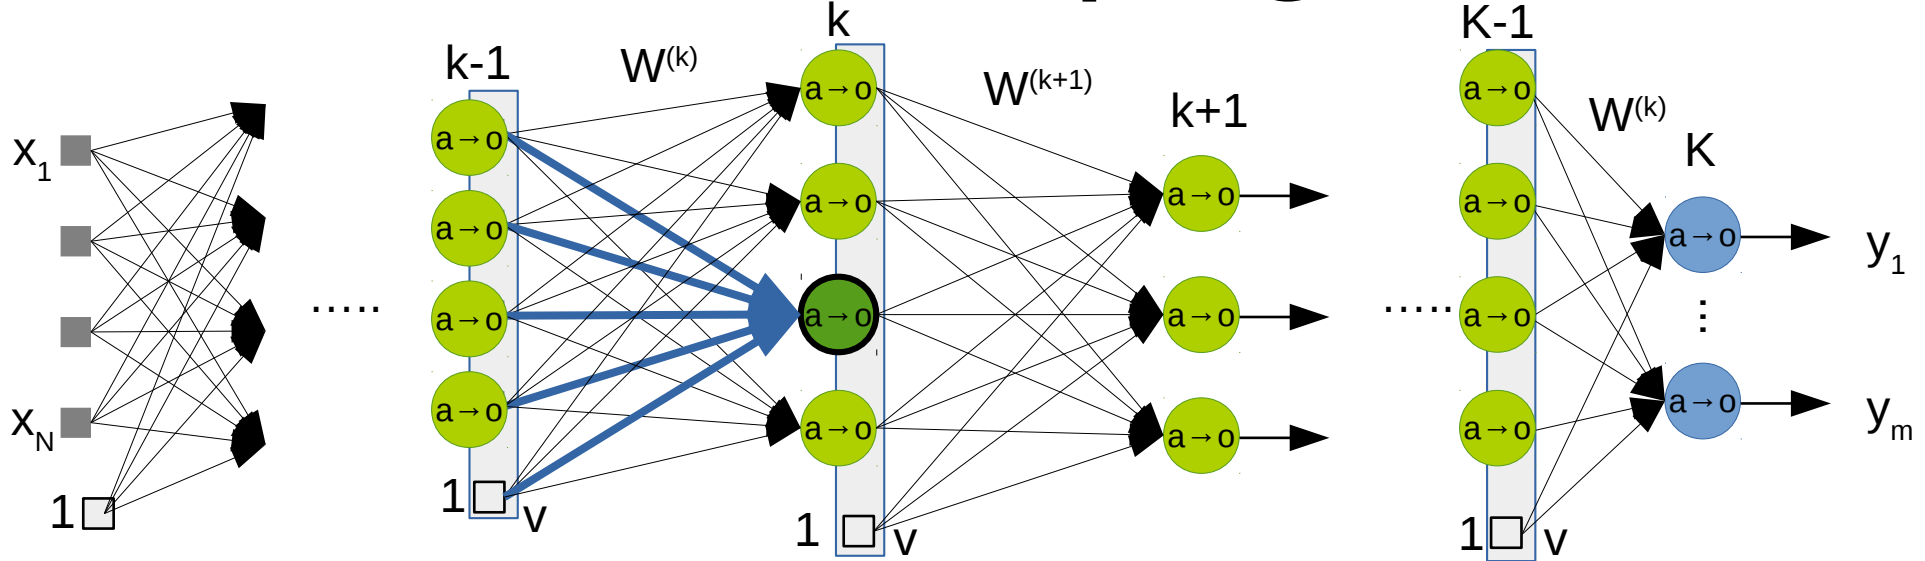


Multi-layer Perceptron (MLP)



- Arbitrary number of hidden layers
- An MLP is a universal approximator
 - Can in principle approximate any continuous function (even with just one hidden layer)
 - Remember: polynomials are universal approximators, too
- Its activations are computed from left to right
 - “Forward propagation” of values
 - As a architecture this is called “feed-forward neural network”

Forward Propagation



Every neuron acts like a perceptron

Multi-layered perceptron with K layers

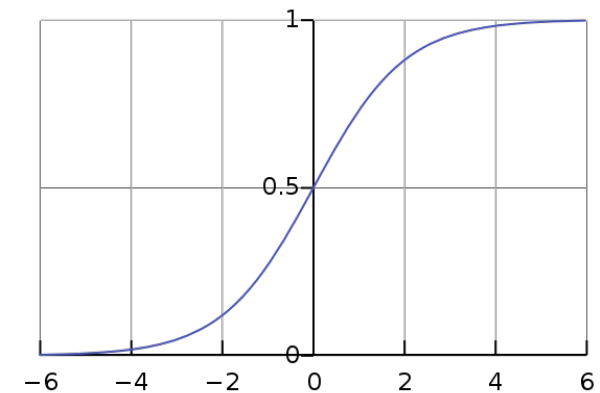
$N^{(k)}$ number of neurons in layer k ($\rightarrow N^{(0)} = n, N^{(K)} = m$)

Synaptic weights from previous neuron j to layer k 's neuron i : $w_{ij}^{(k)}$

Net activation i in layer k : $a_i^{(k)} = \sum_{j=1}^{N^{(k-1)}+1} w_{ij}^{(k)} \cdot v_j^{(k-1)}$

Output activation: $o_i^{(k)} = f(a_i^{(k)})$

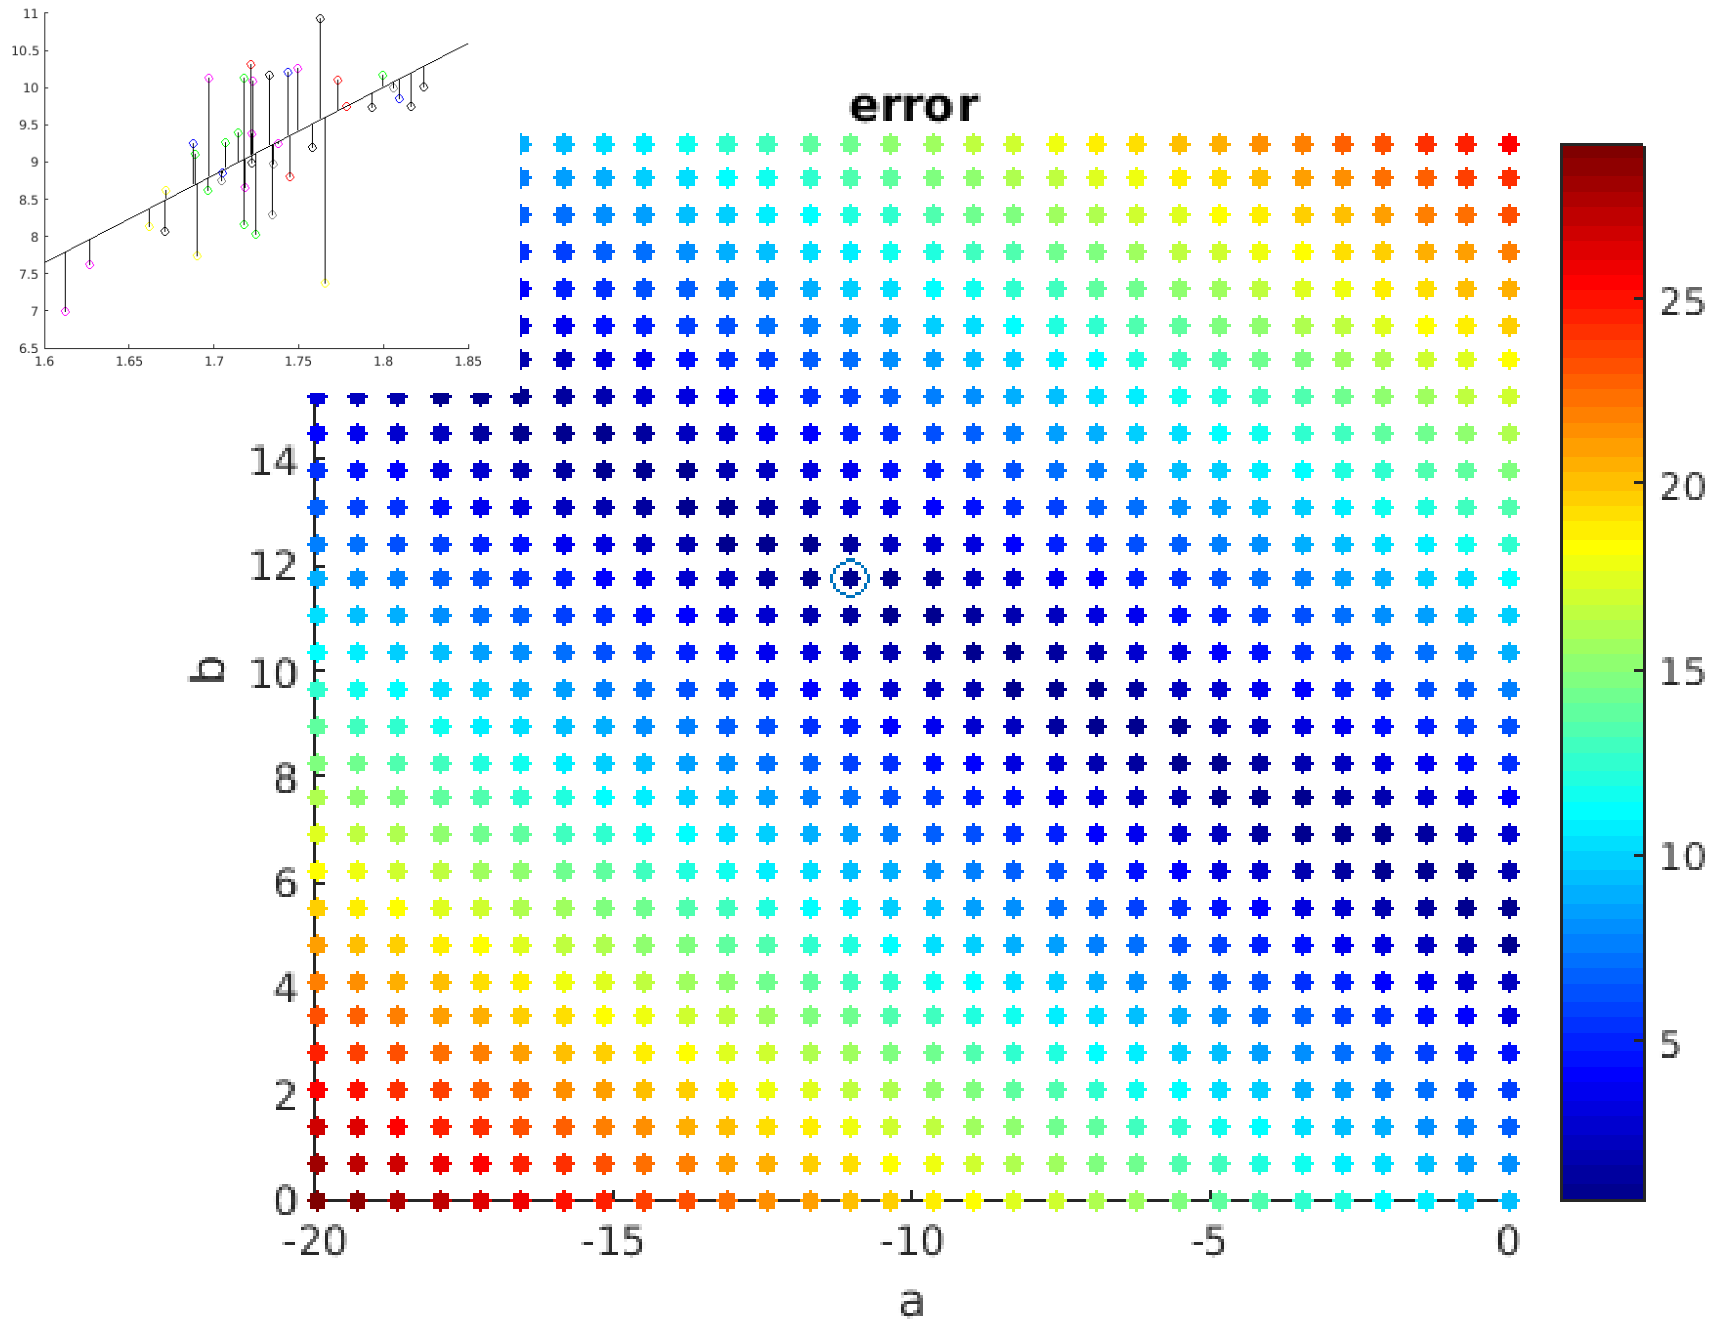
f is the sigmoid function



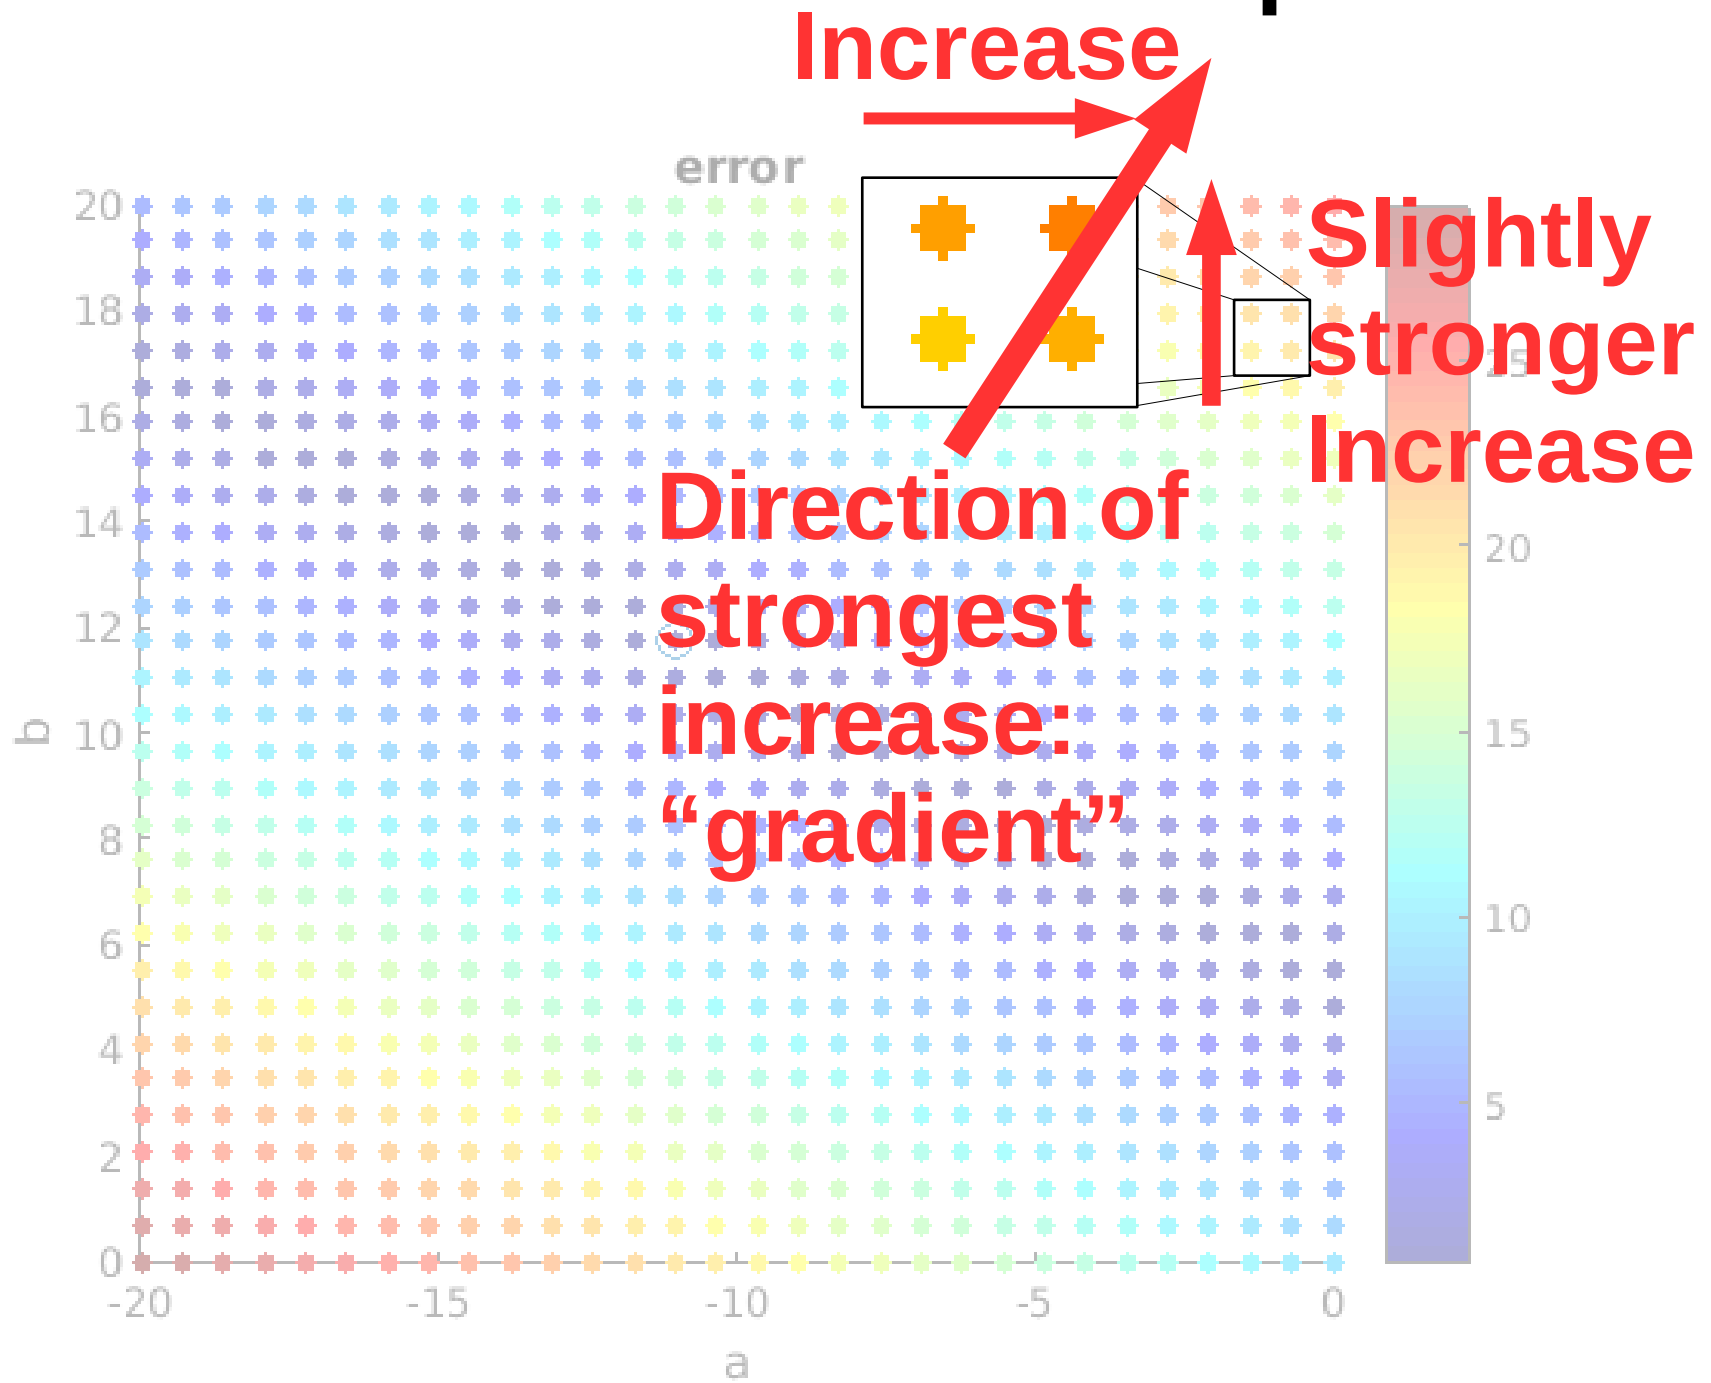
Learning

- How to optimize the parameters?

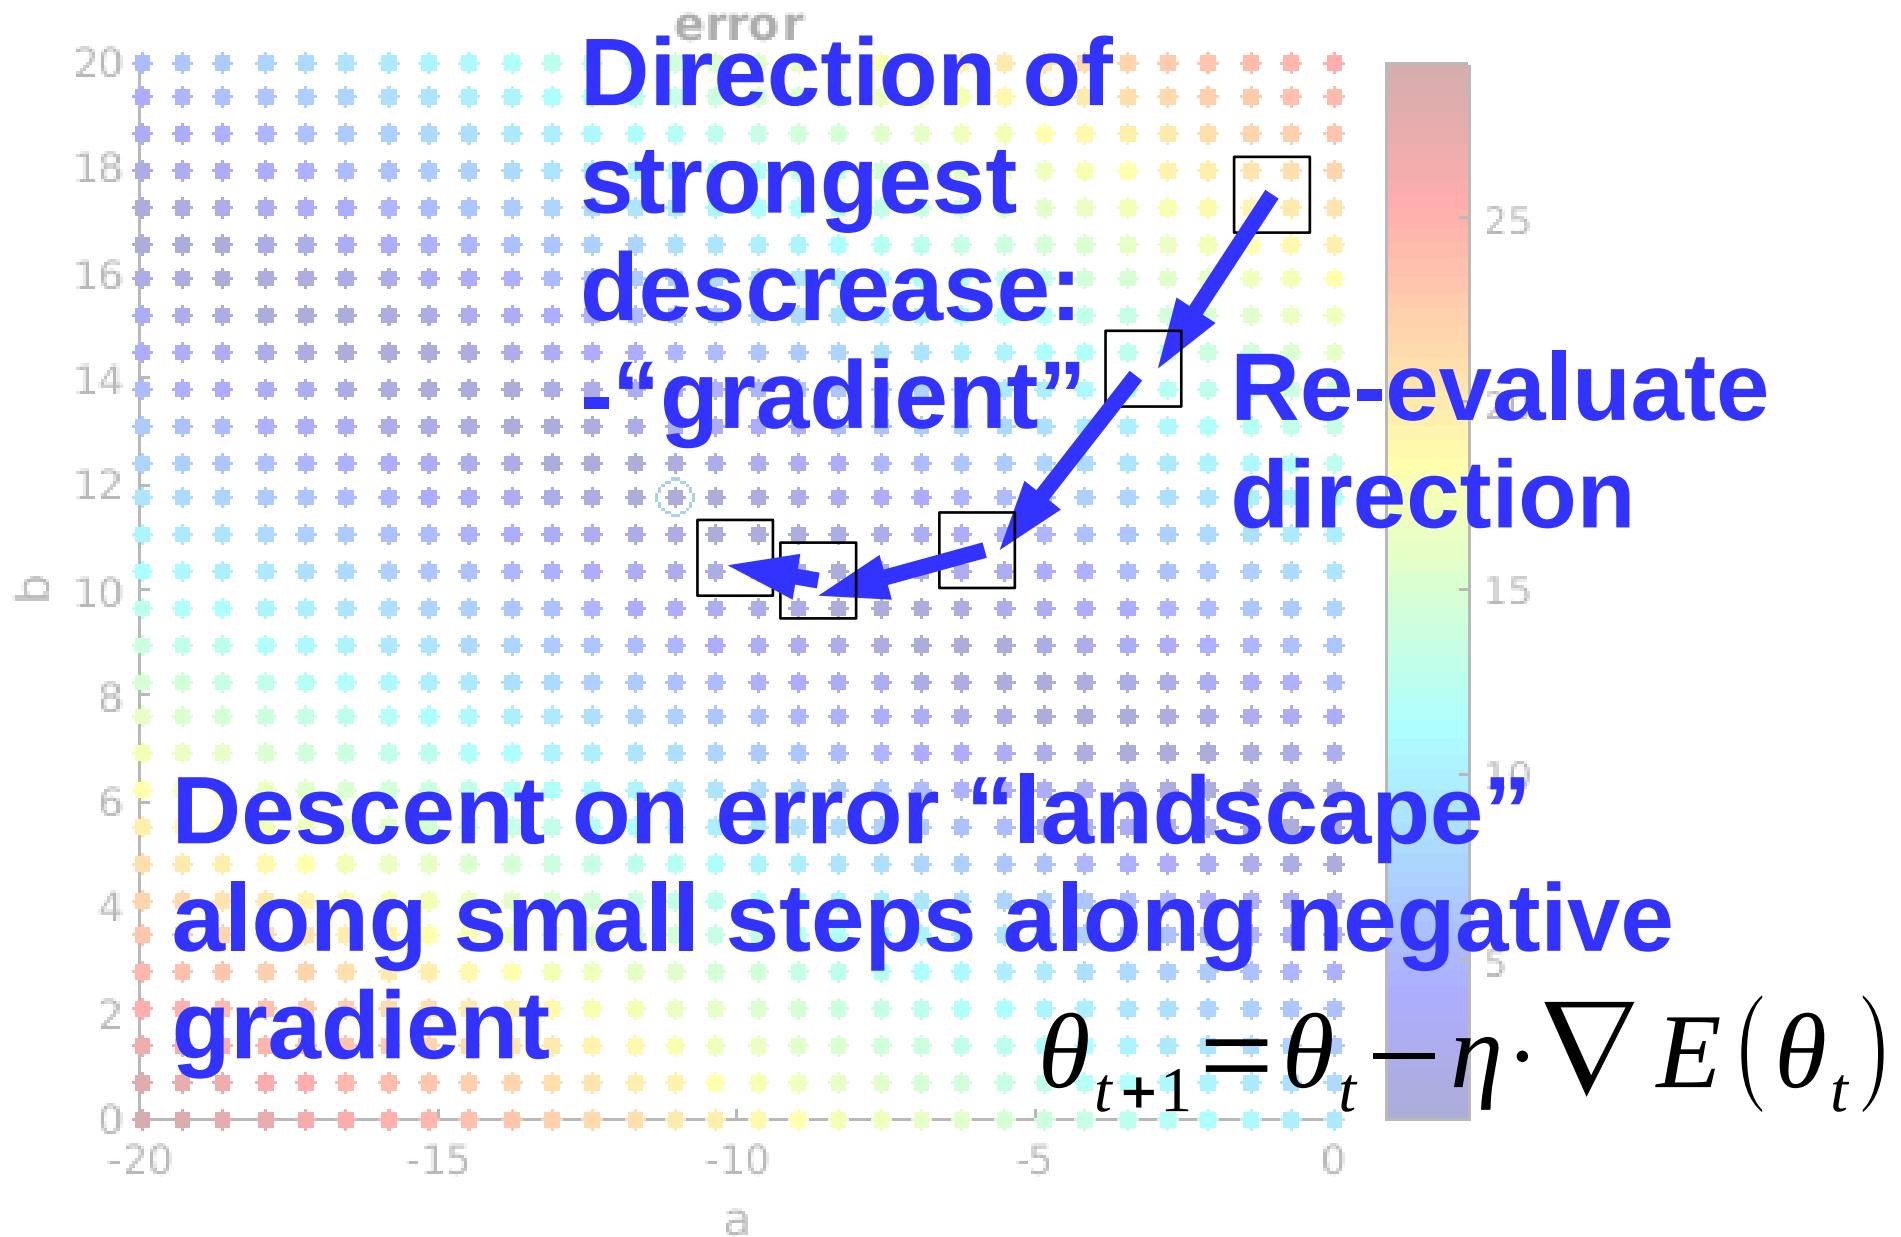
Error Landscape



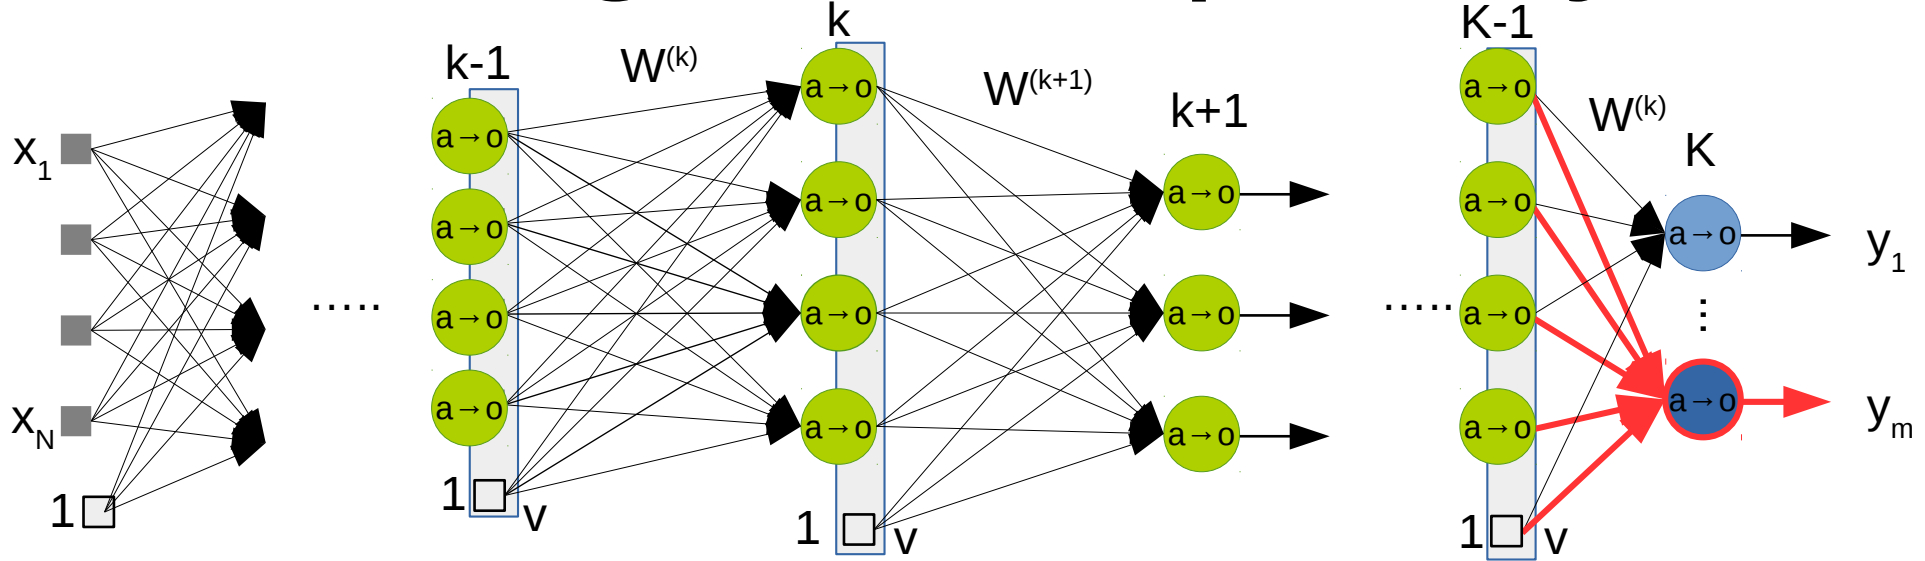
Error Landscape



Error Gradient descent

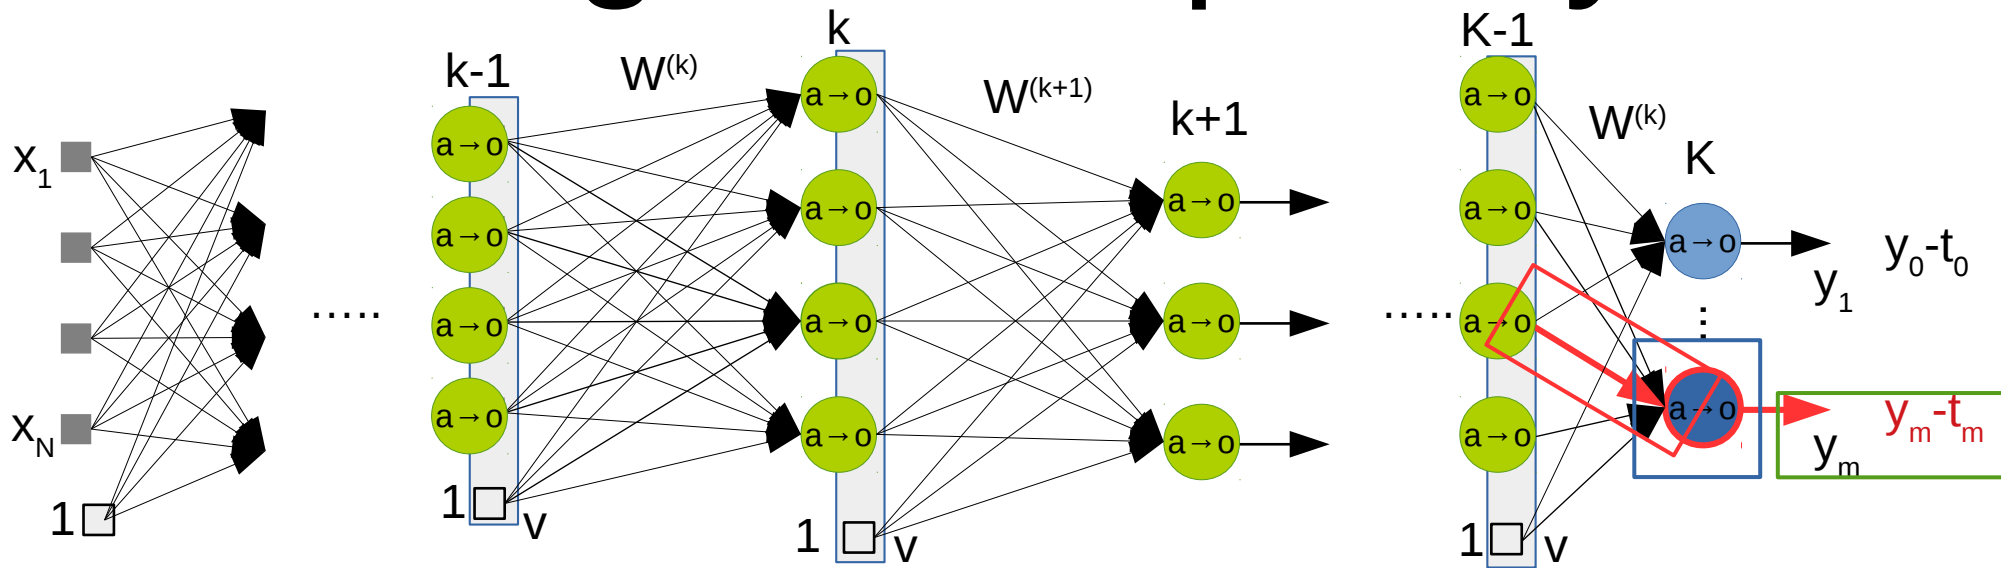


Training the output layer



- The last layer can be trained almost like linear regression
- There is a target output available
- Gradient descent!

Training the output layer



Supervision: target output \vec{t}

$$E = \|\vec{y} - \vec{t}\|^2 = (\vec{y} - \vec{t})^T \cdot (\vec{y} - \vec{t})$$

$$\vec{y} = \vec{o} = f(W^{(K)} \cdot \vec{v}^{(K-1)})$$

Chain rule gives the derivative for a single weight:

$$\begin{aligned} \frac{\partial E}{\partial w_{ij}^{(K)}} &= \frac{\partial E}{\partial o_i^{(K)}} \cdot \frac{\partial o_i^{(K)}}{\partial a_i^{(K)}} \cdot \frac{\partial a_i^{(K)}}{\partial w_{ij}^{(K)}} \\ &= [o_i^{(K)} - t_i] \cdot [o_i^{(K)} \cdot (1 - o_i^{(K)})] \cdot v_j \end{aligned}$$

How much does wiggling w wiggle E ?
=

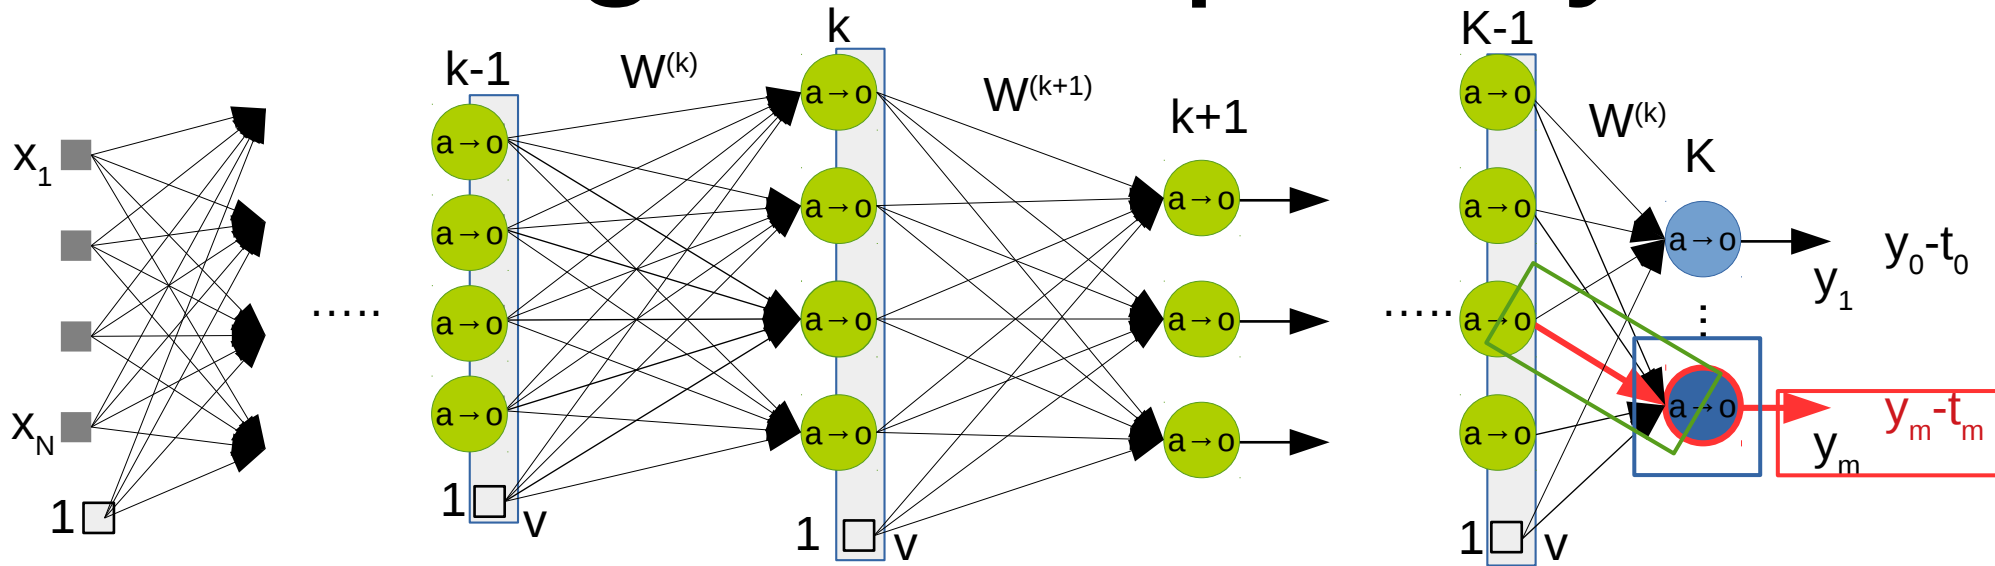
How much does wiggling w wiggle a ?

How much does wiggling a wiggle o ?

How much does wiggling o wiggle E ?

$$\text{Gradient descent update: } w_{ij}^{(K)} \leftarrow w_{ij}^{(K)} - \eta \cdot \frac{\partial E}{\partial w_{ij}^{(K)}}$$

Training the output layer



Supervision: target output \vec{t}

$$E = \|\vec{y} - \vec{t}\|^2 = (\vec{y} - \vec{t})^T \cdot (\vec{y} - \vec{t})$$

$$\vec{y} = \vec{o} = f(W^{(K)} \cdot \vec{v}^{(K-1)})$$

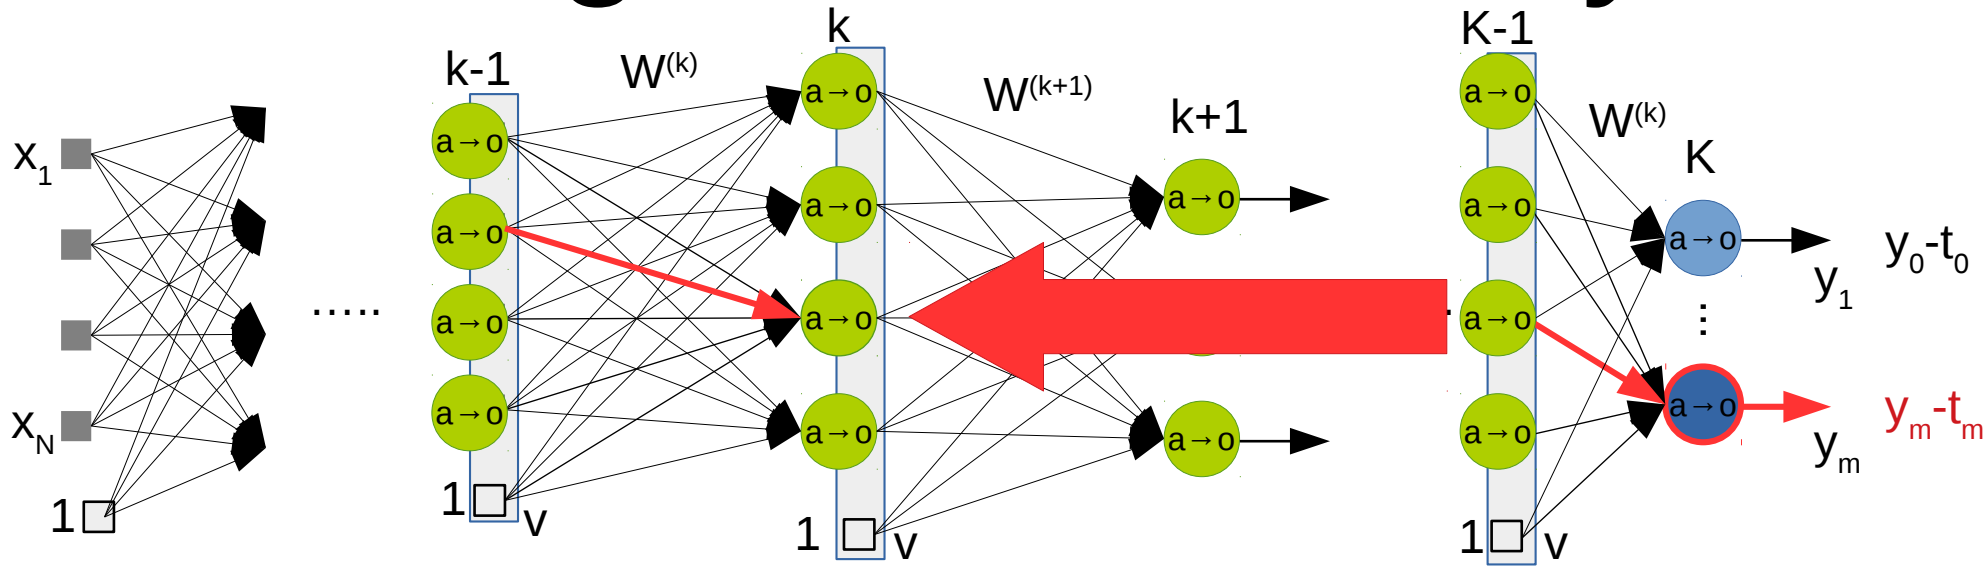
Chain rule gives the derivative for a single weight:

$$\begin{aligned} \frac{\partial E}{\partial w_{ij}^{(K)}} &= \frac{\partial E}{\partial o_i^{(K)}} \cdot \frac{\partial o_i^{(K)}}{\partial a_i^{(K)}} \cdot \frac{\partial a_i^{(K)}}{\partial w_{ij}^{(K)}} \\ &= [o_i^{(K)} - t_i] \cdot [o_i^{(K)} \cdot (1 - o_i^{(K)})] \cdot v_j \end{aligned}$$

**This is like the previous perceptron:
Change = error x input**

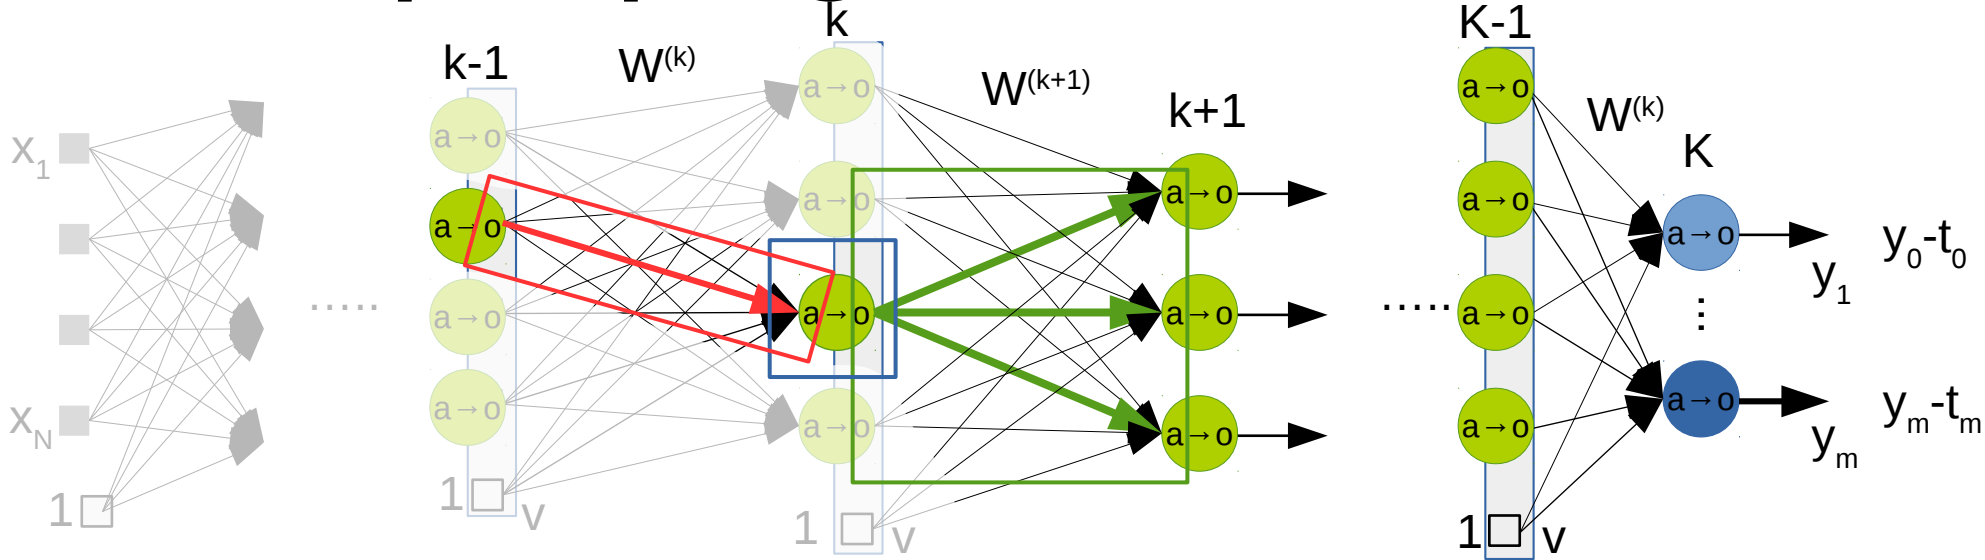
Gradient descent update: $w_{ij}^{(K)} \leftarrow w_{ij}^{(K)} - \eta \cdot \frac{\partial E}{\partial w_{ij}^{(K)}}$

Training the hidden layer



- How to get the gradient for the hidden layers?
- No direct target output for the neurons available!
- Approach: propagate error backwards

Backpropagation of errors



Chain rule gives the derivative for a single weight (same as before):

$$\frac{\partial E}{\partial w_{ij}^{(k)}} = \frac{\partial E}{\partial o_i^{(k)}} \cdot \frac{\partial o_i^{(k)}}{\partial a_i^{(k)}} \cdot \frac{\partial a_i^{(k)}}{\partial w_{ij}^{(k)}}$$

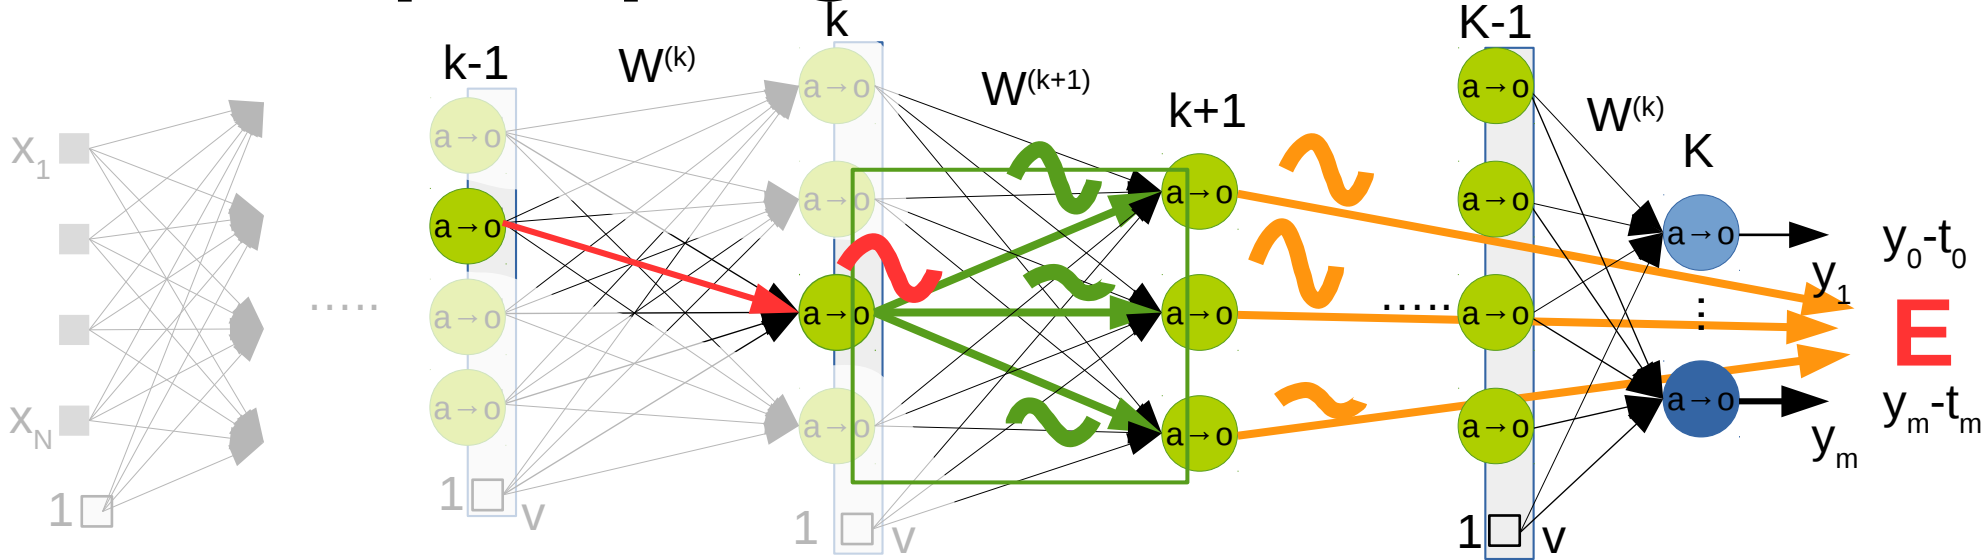
$$= [?] \cdot [o_i^{(k)} \cdot (1 - o_i^{(k)})] \cdot v_j$$

Gradient descent: $w_{ij}^{(k)} \leftarrow w_{ij}^{(k)} - \eta \cdot \frac{\partial E}{\partial w_{ij}^{(k)}}$

The 2nd and 3rd term are exactly the same as for the output layer.

The 1st term has to be determined through recursion through the green weights.

Backpropagation of errors



Chain rule gives the derivative for a single weight (same as before):

$$\frac{\partial E}{\partial w_{ij}^{(k)}} = \frac{\partial E}{\partial o_i^{(k)}} \cdot \frac{\partial o_i^{(k)}}{\partial a_i^{(k)}} \cdot \frac{\partial a_i^{(k)}}{\partial w_{ij}^{(k)}}$$

$$= [?????] \cdot [o_i^{(k)} \cdot (1 - o_i^{(k)})] \cdot v_j$$

Gradient descent: $w_{ij}^{(k)} \leftarrow w_{ij}^{(k)} - \eta \cdot \frac{\partial E}{\partial w_{ij}^{(k)}}$

How much does wiggling o wiggle E?

=

How does wiggling o wiggle a_1 wiggle E?

+

How does wiggling o wiggle a_2 wiggle E?

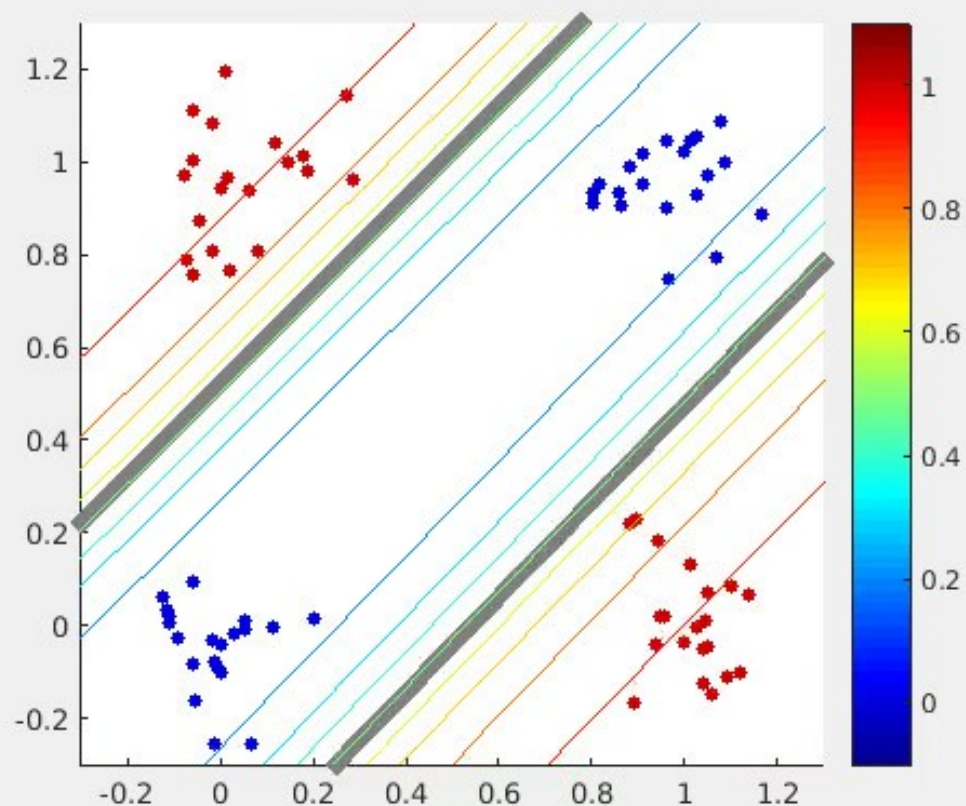
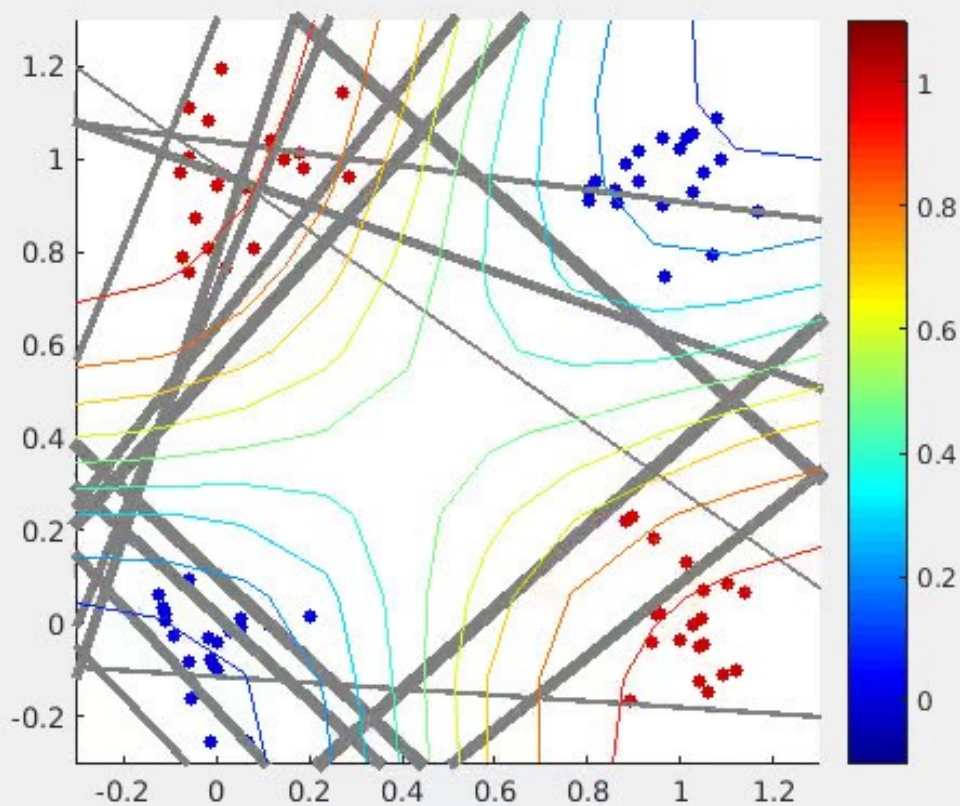
+

...

$$\frac{\partial E}{\partial o_i^{(k)}} = \sum_{l=1}^{N^{(k+1)}} \frac{\partial E}{\partial a_l^{(k+1)}} \cdot \frac{\partial a_l^{(k+1)}}{\partial o_i^{(k)}} = \sum_{l=1}^{N^{(k+1)}} \frac{\partial E}{\partial a_l^{(k+1)}} \cdot w_{li}^{(k)}$$

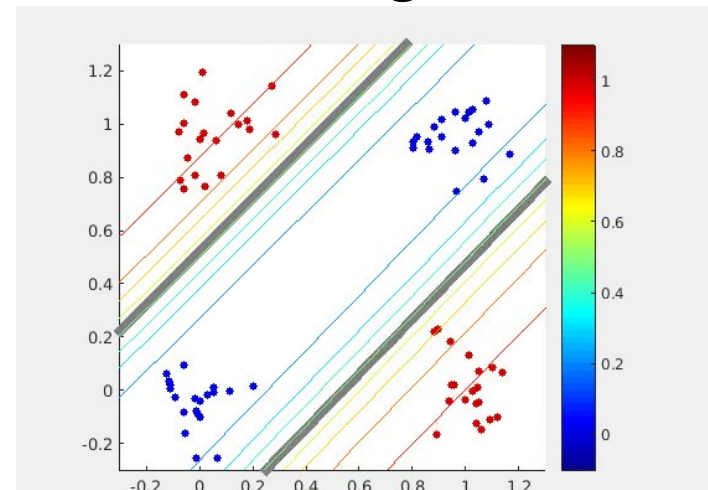
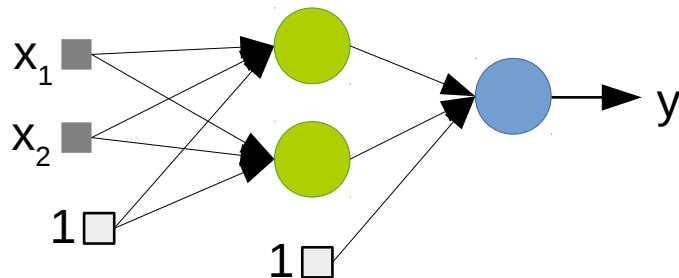
The orange term was computed in the previous layer! Recursion complete!

XOR Revisited



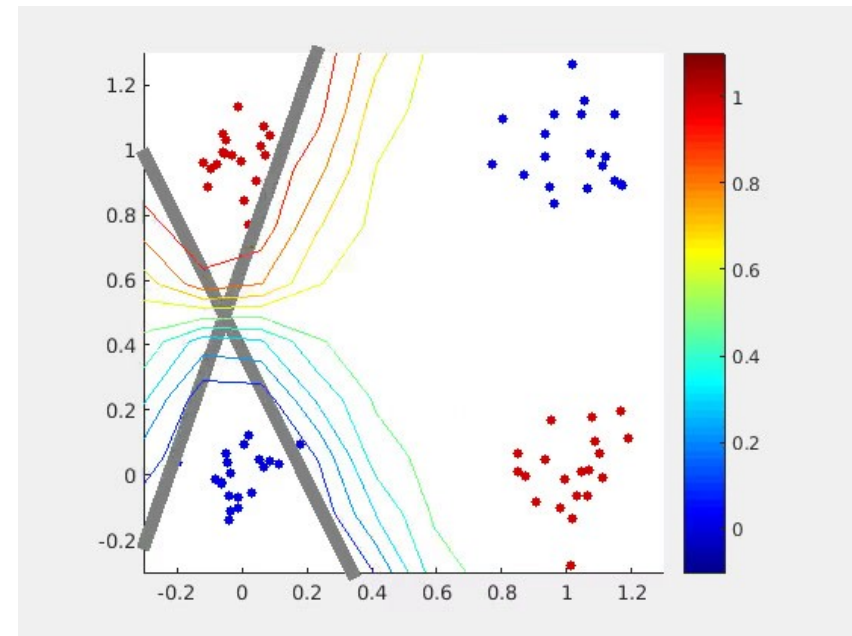
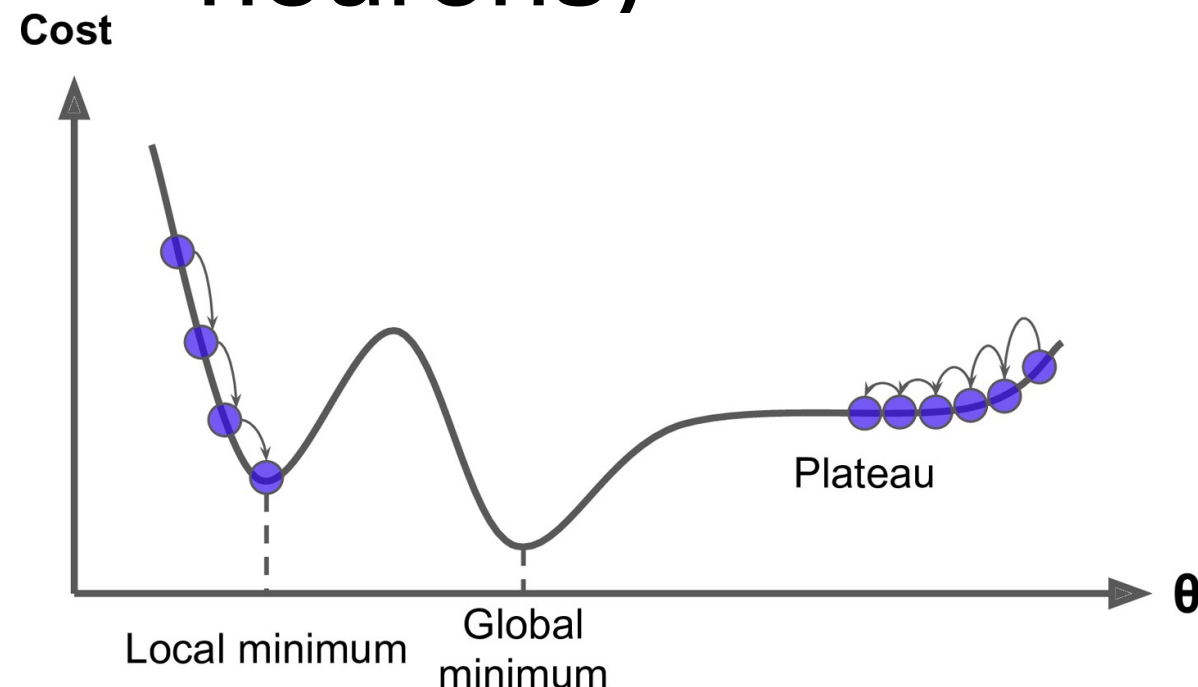
Random Initialization

- Neural network weights cannot be initialized as zero
- Random initialization necessary!
- Same initialization would make all features the same \rightarrow same gradient \rightarrow always the same features
- This makes MLP a stochastic algorithm that needs to be run several times to estimate how good it is!



Local Optima

- MLP do have local optima!
- They also have several equal global optima (by switching neurons)



Practical

- The following code allows to hand-code a neural network by choosing weights

```
library(nnet)
```

```
# bias,x1,x2 -> first hidden, then second hidden; bias,h1,h2 -> output
```

```
W <- c(-3,10,0,-3,0,10,-2,2,2)
```

```
nn <- nnet(matrix(c(0,0),nrow=1, ncol=2), matrix(c(0), nrow=1, ncol=1), size = 2, maxit = 1)
```

```
# the following line hand-sets the weights
```

```
# nnet can be trained automatically by providing real data above and removing the following line
```

```
nn$wts <- W
```

```
x <- (0:20) * 0.05
```

```
y <- x
```

```
xy <- expand.grid(x, y)
```

```
z <- predict(nn, xy)
```

```
z <- matrix(z,nrow=21,ncol=21,byrow=TRUE)
```

```
contour(x,y,z)
```

- Run the code and try to understand the neuron's output
- Choose the weights differently so that the neural network implements a XOR function of the inputs

Practical

- Start preparing for the coursework interim assessment (week 6)
- Inspect loadMNIST.R script uploaded on moodle